PRISMA ORM

CHEAT SHEETS

A Complete Reference for Prisma ORM 7 with TypeScript

Schema · Prisma Client · CRUD · Filtering · Pagination · Relations · Transactions · Migrations · Errors · Performance · $extends · TypeScript · Testing · Raw SQL · Security

34 REFERENCE SHEETS

CONTENTS

INTRODUCTION & SETUP

01 Introduction & Installation
02 Prisma Client Setup & Methods

SCHEMA DEFINITION

03 Schema: Models & Fields
04 Schema: Relations & Enums
05 Schema: Indexes, Constraints & Attributes

CRUD OPERATIONS

06 CRUD: Create Operations
07 CRUD: Read Operations
08 CRUD: Update & Upsert
09 CRUD: Delete & Batch Operations

FILTERING & QUERYING

10 Filtering: Where Clauses & Operators
11 Filtering: Advanced & Pagination

RELATIONS

12 Relations: Include & Select
13 Relations: Nested Writes
14 Relations: Many-to-Many & Advanced

ADVANCED TOPICS

15 Aggregations & Grouping
16 Transactions
17 Migrations
18 Error Handling
19 Performance: N+1 & Indexing
20 Connection Pooling & Optimization

ADVANCED TOPICS (CONT.)

21 Advanced: Soft Deletes & Auditing
22 Advanced: $extends & Custom Methods
23 TypeScript & Type Safety
24 Testing & Mocking
25 Debugging & Studio

SPECIALIZED TOPICS

26 Raw SQL Queries
27 Special Cases: Null, JSON, Arrays
28 Self-Referential & Complex Relations
29 Validation & Custom Logic
30 Best Practices & Patterns
31 Security & Data Protection
32 Migration Strategies

QUICK REFERENCE

33 Quick Reference — Schema & CLI
34 Quick Reference — Client Queries

PRISMA

Table of Contents

34 sheets · syntax · examples · notes

INTRODUCTION & SETUP 01 – 02

01 **Introduction & Installation**
prisma init · prisma.config.ts · DATABASE_URL · driver adapter · prisma generate

02 **Prisma Client Setup & Methods**
new PrismaClient · adapter · singleton · $transaction · $queryRaw · $disconnect

SCHEMA DEFINITION 03 – 05

03 **Schema: Models & Fields**
model · @id · @default · @unique · @updatedAt · ? optional · [] list

04 **Schema: Relations & Enums**
@relation · fields · references · onDelete · enum · implicit m-n

05 **Schema: Indexes, Constraints & Attributes**
@@index · @@unique · @@id · @db.VarChar · sort · type: Gin

CRUD OPERATIONS 06 – 09

06 **CRUD: Create Operations**
create · createMany · createManyAndReturn · skipDuplicates · nested create

07 **CRUD: Read Operations**
findUnique · findFirst · findMany · select · include · omit · orderBy · skip/take

08 **CRUD: Update & Upsert**
update · updateMany · updateManyAndReturn · increment · upsert · nested update

09 **CRUD: Delete & Batch Operations**
delete · deleteMany · onDelete: Cascade · P2003 · soft delete · $transaction

FILTERING & QUERYING 10 – 11

10 **Filtering: Where Clauses & Operators**
equals · gt/gte/lt/lte · contains · startsWith · in/notIn · not · AND/OR/NOT

11 **Filtering: Advanced & Pagination**
cursor · skip/take · some/every/none · is/isNot · has/hasSome · nulls · _count

RELATIONS 12 – 14

12 **Relations: Include & Select**
include · select · nested where/orderBy/take · _count · relationLoadStrategy

13 **Relations: Nested Writes**
create · connect · connectOrCreate · disconnect · set · update · upsert · delete

14 **Relations: Many-to-Many & Advanced**
implicit m-n · explicit join model · @@id · connect · set · some

ADVANCED TOPICS 15 – 25

15 **Aggregations & Grouping**
aggregate · _count · _sum · _avg · _min · _max · groupBy · having

16 **Transactions**
$transaction([…]) · interactive tx · isolationLevel · timeout · maxWait · P2034

17 **Migrations**
migrate dev · db push · migrate deploy · migrate status · migrate resolve · migrate reset

18 **Error Handling**
PrismaClientKnownRequestError · P2002 · P2025 · P2003 · …OrThrow · retry

19 **Performance: N+1 & Indexing**
N+1 · include · in · fluent API batching · relationLoadStrategy · @@index

20 **Connection Pooling & Optimization**
PrismaPg · max · connectionTimeoutMillis · idleTimeoutMillis · PgBouncer · $on('query')

21 **Advanced: Soft Deletes & Auditing**
deletedAt · $extends query · restore · AuditLog · createdAt · @updatedAt

22 **Advanced: $extends & Custom Methods**
$extends · model · query · result · client · $allModels · Prisma.defineExtension

23 **TypeScript & Type Safety**
Prisma.UserGetPayload · satisfies Prisma.UserSelect · UserWhereInput · Prisma.Result

24 **Testing & Mocking**
mockDeep · __mocks__ · vi.mock · .env.test · migrate reset · deleteMany · factories

25 **Debugging & Studio**
prisma studio · log · $on('query') · DEBUG · prisma validate · prisma debug

SPECIALIZED TOPICS 26 – 32

26 **Raw SQL Queries**
$queryRaw · $executeRaw · Prisma.sql · Prisma.join · $queryRawUnsafe · TypedSQL

27 **Special Cases: Null, JSON, Arrays**
null · Prisma.JsonNull · DbNull · path · String[] · BigInt · Bytes · Decimal

28 **Self-Referential & Complex Relations**
self-relation · manager/reports · parentId tree · followers m-n · WITH RECURSIVE

29 **Validation & Custom Logic**
zod · safeParse · z.infer · $extends query · CHECK constraint · business rules

30 **Best Practices & Patterns**
naming · multi-file schema · repository · select shapes · error mapping · project layout

31 **Security & Data Protection**
SQL injection · tenant scoping · RLS · omit · sslmode · least privilege

32 **Migration Strategies**
expand-contract · --create-only · backfill · CONCURRENTLY · rollback · backups

QUICK REFERENCE 33 – 34

33 **Quick Reference — Schema & CLI**
CLI commands · prisma.config.ts · generator & datasource · field types · attributes · relations · enums · migrations workflow

34 **Quick Reference — Client Queries**
CRUD methods · filters & operators · relation queries · nested writes · pagination · aggregation · transactions · raw SQL · errors

Introduction & Installation

WHAT IS PRISMA ORM?

Prisma ORM is a schema-first, type-safe ORM for Node.js and TypeScript. You model data in `schema.prisma`, **Prisma Migrate** turns it into tables, and **Prisma Client** is generated from it so every query is type-checked. Prisma 7 is Rust-free: the client connects through a JavaScript **driver adapter**, and the connection URL lives in `prisma.config.ts`.

REQUIREMENTS (V7)

- Node ≥ 20.19 · TypeScript ≥ 5.4
- ESM: `"type": "module"` in package.json
- `.env` is not auto-loaded — use `dotenv/config`
- Postgres, MySQL/MariaDB, SQLite, SQL Server, CockroachDB

NEW PROJECT (POSTGRESQL)

```
npm init -y && npm pkg set type=module
npm i -D prisma typescript tsx @types/node
npm i @prisma/client @prisma/adapter-pg dotenv
npx tsc --init
npx prisma init --datasource-provider postgresql \
  --output ../generated/prisma
# creates prisma/schema.prisma, prisma.config.ts, .env
# .env
DATABASE_URL="postgresql://app:secret@localhost:5432/shop"
npx prisma migrate dev --name init  # SQL migration + apply
npx prisma generate                 # build the typed client
npx tsx src/index.ts                # run your code
```

CODE EXPLAINED

- → `prisma` — CLI: init, migrate, generate, studio
- → `@prisma/client` — runtime for the generated client
- → `@prisma/adapter-pg` — driver adapter (pg)
- → `--output` — client location; import from there
- → `migrate dev` — schema diff → SQL → apply
- → `generate` — re-run after every schema change (v7 no longer runs it from `migrate dev`)

PRISMA.CONFIG.TS + SCHEMA.PRISMA

```
// prisma.config.ts (project root)
import "dotenv/config";
import { defineConfig, env } from "prisma/config";
export default defineConfig({
  schema: "prisma/schema.prisma",
  datasource: { url: env("DATABASE_URL") },
});
// prisma/schema.prisma
generator client {
  provider = "prisma-client"
  output   = "../generated/prisma"
}
datasource db {
  provider = "postgresql"
}
```

DATABASE_URL FORMATS

PROVIDER	EXAMPLE
postgresql	postgresql://u:p@host:5432/db
mysql	mysql://u:p@host:3306/db
sqlite	file:./dev.db
cockroachdb	postgresql://u:p@host:26257/db

DRIVER ADAPTERS

DATABASE	PACKAGE
PostgreSQL	@prisma/adapter-pg
MySQL/MariaDB	@prisma/adapter-mariadb
SQLite	@prisma/adapter-better-sqlite3
SQL Server	@prisma/adapter-mssql

⚠ COMMON MISTAKES

- Importing from `@prisma/client` — import from your `output` path
- Keeping `url` in `datasource` — v7 rejects it; use the config file
- Skipping `prisma generate` after schema edits — stale types
- Committing `.env` — add it to `.gitignore`

✓ TIPS & BEST PRACTICES

- Add `generated/` to `.gitignore`; generate in `postinstall`
- Install the VS Code Prisma extension for autocomplete
- `npx prisma format` fixes layout and missing relations
- `npx prisma init --db` provisions a hosted Prisma Postgres

QUICK SUMMARY

- `prisma init` — schema, config, .env
- `prisma.config.ts` — paths + datasource URL
- `prisma-client` — v7 generator, needs `output`
- `prisma generate` — rebuild client

Prisma Client Setup & Methods

CREATING THE CLIENT

`PrismaClient` is the generated query builder. Each instance owns a **connection pool**, so create **one instance per process** and import it everywhere. Model delegates (`prisma.user`) expose CRUD methods; $-prefixed methods act on the client itself: transactions, raw SQL, extensions, and lifecycle.

NOTES

- Queries run only when awaited
- Pool connects lazily — `$connect()` optional
- `model BlogPost` → `prisma.blogPost`
- Dev hot-reload → cache on `globalThis`

SRC/LIB/PRISMA.TS — SINGLETON

```ts
import "dotenv/config";
import { PrismaClient } from "../../generated/prisma/client";
import { PrismaPg } from "@prisma/adapter-pg";

const adapter = new PrismaPg({ connectionString: process.env.DATABASE_URL! });
const create = () => new PrismaClient({ adapter, log: ["warn", "error"] });

const g = globalThis as unknown as { prisma?: ReturnType<typeof create> };
export const prisma = g.prisma ?? create();
if (process.env.NODE_ENV !== "production") g.prisma = prisma;
// elsewhere
import { prisma } from "./lib/prisma";
const user = await prisma.user.findUnique({ where: { id: 1 } }); // User | null

// graceful shutdown
process.on("SIGTERM", async () => {
  await prisma.$disconnect();
  process.exit(0);
});
```

CODE EXPLAINED

- → `PrismaPg` — wraps the `pg` driver + pool
- → `log` — `query` · `info` · `warn` · `error`
- → `globalThis` — one pool across hot reloads
- → `findUnique` — result type inferred
- → `$disconnect()` — close pool on shutdown
- → Serverless: create outside the handler

CLIENT ($) METHODS

METHOD	PURPOSE
`$transaction([q1, q2])`	Atomic batch
`$transaction(async tx => …)`	Interactive tx
`$queryRaw`…``	Raw SELECT → rows
`$executeRaw`…``	Raw write → count
`$queryRawUnsafe(sql)`	String SQL (risky)
`$extends({...})`	Client extensions
`$on('query', fn)`	Log events
`$connect / $disconnect`	Open / close pool

CONSTRUCTOR OPTIONS

OPTION	EXAMPLE
`adapter`	`new PrismaPg({...})` — required
`log`	`['query', 'error']`
`errorFormat`	`'pretty' · 'minimal'`
`omit`	`{ user: { password: true } }`
`transactionOptions`	`{ timeout: 10000 }`

⚠ COMMON MISTAKES

- `new PrismaClient()` per request — exhausts connections
- Missing `await` — the query never runs
- `$disconnect()` after each request — slow reconnects
- No `dotenv/config` — `DATABASE_URL` is undefined

✓ TIPS & BEST PRACTICES

- Export one `prisma` from `lib/prisma.ts`
- Global `omit` hides `password` everywhere
- Use `…OrThrow` when a missing row is a bug
- Set `transactionOptions` once, not per call

QUICK SUMMARY

- `PrismaClient({ adapter })` — one per process
- `prisma.model.method()` — typed CRUD
- `$transaction` — batch or callback
- `$disconnect()` — close on shutdown

Schema: Models & Fields

MODELS

A `model` maps to a table; each line is a field: `name Type modifier @attributes`. Scalar types map to both a database column and a TypeScript type. Every model needs a unique identifier (`@id` or `@@id`). `@` attributes configure one field; `@@` block attributes configure the whole model.

CONVENTIONS

- Models PascalCase singular: `OrderItem`
- Fields camelCase: `createdAt`
- `@@map("users")` for snake_case tables
- `prisma format` aligns and fixes

SCHEMA.PRISMA

```
model User {
  id         Int       @id @default(autoincrement())
  publicId   String    @unique @default(uuid(7))
  email      String    @unique
  name       String?                  // string | null
  isActive   Boolean   @default(true)
  balance    Decimal   @default(0) @db.Decimal(12, 2)
  tags       String[]                 // Postgres list
  settings   Json      @default("{}")
  views      BigInt    @default(0)
  createdAt  DateTime  @default(now())
  updatedAt  DateTime  @updatedAt
  lastLogin  DateTime? @map("last_login")

  @@map("users")
}
```

CODE EXPLAINED

- → `autoincrement()` — DB sequence PK
- → `uuid(7)` — time-sortable UUID
- → `String?` — nullable column
- → `@db.Decimal(12, 2)` — native type
- → `String[]` — not on MySQL/SQLite
- → `@updatedAt` — set by Prisma on update
- → `@map` / `@@map` — DB-side names

SCALAR TYPES

PRISMA	TS	POSTGRES
`String`	`string`	`text`
`Boolean`	`boolean`	`boolean`
`Int`	`number`	`integer`
`BigInt`	`bigint`	`bigint`
`Float`	`number`	`float8`
`Decimal`	`Decimal`	`numeric`
`DateTime`	`Date`	`timestamp(3)`
`Json`	`JsonValue`	`jsonb`
`Bytes`	`Uint8Array`	`bytea`

@DEFAULT() VALUES

FUNCTION	GENERATES
`autoincrement()`	DB sequence
`uuid()` / `uuid(7)`	UUID v4 / v7
`cuid(2)`	CUID v2
`nanoid(16)`	NanoID
`ulid()`	ULID
`now()`	Timestamp
`dbgenerated("…")`	SQL expression
`true` · `0` · `"x"` · `[]`	Literal

⚠ COMMON MISTAKES

- `Float` for money — rounding errors; use `Decimal`
- New required field on a full table — needs a `@default`
- Expecting `@updatedAt` to change on raw SQL updates
- `+` on `Decimal` — use `.plus()` / `.times()`

✓ TIPS & BEST PRACTICES

- Public IDs: `uuid(7)` / `cuid(2)`; joins: `autoincrement()`
- Store `DateTime` in UTC; format in the UI
- `@db.VarChar(n)` for a DB-enforced length
- `@ignore` hides a column from the client

QUICK SUMMARY

- `model` — one table, one delegate
- `?` / `[]` — optional / list
- `@id @default()` — PK + generator
- `@updatedAt` — auto timestamp

Schema: Relations & Enums

RELATIONS & ENUMS

A relation pairs a **relation field** (Prisma-only) with a **foreign key** column. The FK side declares `@relation(fields, references)`; the other side holds a back-relation (`Post[]` or `Profile?`). Implicit many-to-many needs only list fields on both sides. An `enum` limits values and becomes a TS union.

NOTES

- 1-1: FK field must be `@unique`
- Two links to one model need names
- `prisma format` adds back-relations
- Enums: native on Postgres/MySQL

SCHEMA.PRISMA

```prisma
enum Role {
  USER
  ADMIN
}

model User {
  id      Int      @id @default(autoincrement())
  role    Role     @default(USER)
  posts   Post[]                          // 1-n
  profile Profile?                        // 1-1
}
model Profile {
  id     Int  @id @default(autoincrement())
  userId Int  @unique
  user   User @relation(fields: [userId], references: [id], onDelete: Cascade)
}
model Post {
  id       Int   @id @default(autoincrement())
  authorId Int?
  author   User? @relation(fields: [authorId], references: [id], onDelete: SetNull)
  tags     Tag[]                // implicit m-n ↔ Tag.posts Post[]
}
```

CODE EXPLAINED

- → `Post[]` — virtual, no column
- → `fields` — FK column(s) here
- → `references` — key on the other model
- → `Cascade` — delete profile with user
- → `SetNull` — keep post, null author
- → `Tag[]` ↔ `Post[]` — join table `_PostToTag`

REFERENTIAL ACTIONS

ACTION	ON PARENT DELETE
`Cascade`	Delete children too
`Restrict`	Block if children exist
`NoAction`	Like Restrict (deferred)
`SetNull`	FK → NULL (optional FK)
`SetDefault`	FK → `@default`
`default (required FK)`	Restrict
`default (optional FK)`	SetNull

RELATION KINDS

KIND	SHAPE
`1-1`	`Profile?` ↔ `userId @unique`
`1-n`	`Post[]` ↔ `authorId`
`m-n implicit`	`Tag[]` ↔ `Post[]`
`m-n explicit`	join model (Sheet 14)
`self`	named relation (Sheet 28)
`enum use`	`Role.ADMIN` in TS

⚠ COMMON MISTAKES

- `@relation(fields…)` on both sides — FK side only
- `SetNull` on a required FK — validation error
- No index on FKs — Postgres doesn't add one
- Comparing enums to raw strings — use `Role.ADMIN`

✓ TIPS & BEST PRACTICES

- Add `@@index([authorId])` on FK columns
- Use explicit m-n once the link needs extra data
- `relationMode = "prisma"` only for DBs without FKs
- Rename enum values in steps: add, migrate, remove

QUICK SUMMARY

- `@relation(fields, references)` — FK side
- `Post[]` / `Profile?` — back-relations
- `onDelete` — Cascade · Restrict · SetNull
- `enum` — fixed values → TS union

Sheet 05

Schema: Indexes, Constraints & Attributes

@@index · @@unique · @@id · @db.VarChar · sort · type: Gin

INDEXES & CONSTRAINTS

Indexes speed up `where`, `orderBy`, and joins but slow writes. Use `@@index` for plain indexes, `@@unique` for unique constraints, and `@@id` for composite keys. Composite indexes follow the **leftmost-prefix rule**: `[tenantId, createdAt]` serves `tenantId` alone or both, not `createdAt` alone.

NOTES

- Unique fields become `findUnique` keys
- Compound key name: `tenantId_slug`
- `name:` renames key · `map:` DB name
- Ship index changes via migrations

SCHEMA.PRISMA

```
model Article {
  id        Int      @id @default(autoincrement())
  tenantId  Int
  slug      String   @db.VarChar(120)
  body      String   @db.Text
  status    Status   @default(DRAFT)
  meta      Json     @db.JsonB
  createdAt DateTime @default(now()) @db.Timestamptz(6)

  @@unique([tenantId, slug])
  @@index([tenantId, status, createdAt(sort: Desc)])
  @@index([meta], type: Gin)
}
model Enrollment {
  studentId Int
  courseId  Int
  @@id([studentId, courseId])
}
// findUnique({ where: { tenantId_slug: { tenantId: 1, slug: "hi" } } })
```

CODE EXPLAINED

- → `@db.VarChar(120)` — DB length limit
- → `@@unique([...])` — unique per tenant
- → `sort: Desc` — matches newest-first sorts
- → `type: Gin` — JSON/array containment
- → `@@id` — composite PK
- → `tenantId_slug` — compound key input

BLOCK ATTRIBUTES

ATTRIBUTE	PURPOSE
`@@id([a, b])`	Composite PK
`@@unique([a, b])`	Composite unique
`@@index([a, b])`	Index
`@@map("t")`	Table name
`@@schema("auth")`	Postgres schema
`@@fulltext([a])`	MySQL full-text

NATIVE TYPES (PG)

ATTRIBUTE	COLUMN
`@db.VarChar(n)`	`varchar(n)`
`@db.Text`	`text`
`@db.Uuid`	`uuid`
`@db.Timestamptz(6)`	`timestamptz`
`@db.Decimal(p,s)`	`numeric(p,s)`
`@db.JsonB`	`jsonb`

WHAT TO INDEX

QUERY	INDEX
`where email`	`@unique`
`where authorId`	`@@index`
`where + orderBy`	both, in order
`multi-tenant`	`tenantId` first
`JSON / array`	`type: Gin`
`boolean flag`	usually skip

⚠ COMMON MISTAKES

- Indexing every column — slower writes, wasted disk
- Wrong composite order — `[createdAt, tenantId]` won't serve `tenantId`
- Shrinking `VarChar(n)` — fails if data is longer
- Missing FK indexes — cascades scan whole tables

✓ TIPS & BEST PRACTICES

- Confirm usage with `EXPLAIN ANALYZE` via `$queryRaw`
- Use `map:` to keep index names stable
- Match index sort order to your hot `orderBy`
- `@@unique` on join models blocks duplicate links

QUICK SUMMARY

- `@@index` — faster filters and sorts
- `@@unique` — uniqueness + `a_b` key
- `@@id` — composite primary key
- `@db.*` — exact native type

CRUD: Create Operations

INSERTING DATA

`create()` inserts one record and returns it; `select`/`include` shape the result. Nested `create` inside `data` writes related rows in the **same transaction**. `createMany()` sends one multi-row INSERT and returns `{ count }`; `createManyAndReturn()` returns the rows (Postgres, CockroachDB, SQLite).

NOTES

- Required fields without defaults must be set
- `createMany` can't nest relations
- `skipDuplicates` → `ON CONFLICT DO NOTHING`
- Duplicate unique value → `P2002`

CREATE EXAMPLES

```ts
const user = await prisma.user.create({
  data: { email: "ada@example.com", name: "Ada" },
});                                     // { id: 1, email: "ada@…", … }

await prisma.user.create({              // nested create, one transaction
  data: {
    email: "cy@example.com",
    profile: { create: { bio: "Hi" } },
    posts: { create: [{ title: "First" }, { title: "Second" }] },
  },
  include: { posts: true },
});

await prisma.post.createMany({
  data: [{ title: "A", authorId: 1 }, { title: "B", authorId: 1 }],
  skipDuplicates: true,
});                                     // { count: 2 }

await prisma.tag.createManyAndReturn({
  data: [{ name: "orm" }, { name: "sql" }],
});                                     // [{ id: 1, name: "orm" }, …]
```

CODE EXPLAINED

- → `data` — typed `UserCreateInput`
- → `profile: { create }` — 1-1 nested insert
- → `posts: { create: [] }` — 1-n nested
- → `include` — return relations too
- → `createMany` — one INSERT, `{ count }`
- → `skipDuplicates` — ignore conflicts
- → `createManyAndReturn` — rows back

CREATE METHODS

METHOD	RETURNS	NESTED
`create`	record	✔
`createMany`	`{ count }`	✖
`createManyAndReturn`	record[]	✖
`upsert`	record	✔
`connectOrCreate`	(nested only)	✔

1,000 ROWS

APPROACH	COST
`for … await create()`	1,000 sequential queries
`Promise.all(creates)`	1,000 queries, pool-bound
`$transaction([...])`	1,000 queries, 1 tx
`createMany({ data })`	1 INSERT — fastest

⚠ COMMON MISTAKES

- Looping `await create()` for bulk data — use `createMany`
- Expecting IDs from `createMany` — use `createManyAndReturn`
- Setting `authorId` and `author: { connect }` together
- Date strings without a timezone — use `Date` or ISO `Z`

✓ TIPS & BEST PRACTICES

- Nested `create` keeps parent + children atomic
- Chunk huge batches (~5,000 rows) to stay under param limits
- Map `P2002` to HTTP 409 Conflict
- `select` only the fields the caller needs

QUICK SUMMARY

- `create` — insert one
- `createMany` — bulk, `{ count }`
- `createManyAndReturn` — bulk + rows
- nested `create` — atomic graph insert

CRUD: Read Operations

QUERYING RECORDS

`findUnique` fetches one record by an `@id`/`@unique` field. `findFirst` returns the first match of any filter. `findMany` returns an array — empty, never `null`. All accept `select`, `include`, or `omit`, and the result type is inferred from what you ask for.

NOTES

- `…OrThrow` variants throw `P2025`
- `select` + `include` can't be siblings
- No `orderBy` = unspecified order
- `findUnique` calls are auto-batched

READ EXAMPLES

```ts
await prisma.user.findUnique({ where: { email: "ada@example.com" } }); // User | null
await prisma.user.findUniqueOrThrow({ where: { id: 1 } });            // or P2025

await prisma.post.findFirst({
  where: { published: true },
  orderBy: { createdAt: "desc" },                 // newest published
});

await prisma.post.findMany({
  where: { authorId: 1, title: { contains: "orm", mode: "insensitive" } },
  select: { id: true, title: true, author: { select: { name: true } } },
  orderBy: [{ createdAt: "desc" }, { id: "desc" }],
  skip: 20,
  take: 10,                                       // page 3
});

await prisma.user.findMany({ omit: { password: true } });
await prisma.user.findMany({ include: { posts: true } });
await prisma.post.findMany({ distinct: ["authorId"] });
await prisma.post.count({ where: { published: true } });   // 42
```

CODE EXPLAINED

- → `findUnique` — by unique key
- → `orderBy` — makes `findFirst` deterministic
- → `mode: 'insensitive'` — Postgres/Mongo
- → nested `select` — shaped payload
- → `id` tie-breaker — stable pages
- → `skip` / `take` — offset paging
- → `omit` — drop secrets
- → `distinct` — unique by field

FINDER COMPARISON

METHOD	WHERE	RESULT
`findUnique`	unique	`T | null`
`findUniqueOrThrow`	unique	T / P2025
`findFirst`	any	`T | null`
`findFirstOrThrow`	any	T / P2025
`findMany`	any	`T[]`
`count`	any	`number`

QUERY ARGUMENTS

ARG	PURPOSE
`where`	Filter (Sheets 10–11)
`select / include`	Pick fields / add relations
`omit`	Drop scalars
`orderBy`	`{ f: 'asc' | 'desc' }`
`skip / take`	Offset / limit
`cursor`	Keyset start

⚠ COMMON MISTAKES

- `findUnique` on a non-unique field — use `findFirst`
- `if (!rows)` on `findMany` — `[]` is truthy; check `.length`
- Huge `skip` values — DB still scans; use a cursor
- `select` and `include` at the same level — error

✓ TIPS & BEST PRACTICES

- Prefer `select` in APIs — small payloads, no leaks
- Add `id` as a final `orderBy` when paginating
- Map `P2025` from `…OrThrow` to HTTP 404
- `count()` beats `findMany().length`

QUICK SUMMARY

- `findUnique` — one by unique key
- `findFirst` — first match
- `findMany` — list, never null
- `select` · `include` · `omit` — shape

CRUD: Update & Upsert

update · updateMany · updateManyAndReturn · increment · upsert · nested update

MODIFYING RECORDS

`update()` changes one record by unique `where` and throws `P2025` if none matches. `updateMany()` changes every match and returns `{ count }`. **Atomic operations** (`increment`, `multiply`…) run in SQL, so concurrent writes never lose updates. `upsert()` updates if the key exists, else creates.

NOTES

- Omitted fields stay unchanged
- `undefined` = skip · `null` = clear
- `@updatedAt` refreshes automatically
- Upsert uses native `ON CONFLICT` when possible

UPDATE EXAMPLES

```ts
await prisma.user.update({
  where: { id: 1 },
  data: { name: "Ada Lovelace", age: null },  // age cleared
});
await prisma.post.update({
  where: { id: 7 },
  data: {
    views: { increment: 1 },                  // views = views + 1
    tags:  { push: "featured" },              // Postgres list append
  },
});
await prisma.post.updateMany({
  where: { authorId: 1, published: false },
  data: { published: true },
});                                           // { count: 3 }
await prisma.user.upsert({
  where:  { email: "ada@example.com" },
  update: { lastLogin: new Date() },
  create: { email: "ada@example.com", lastLogin: new Date() },
});
```

CODE EXPLAINED

- → `where` — unique field required
- → `null` — set column to NULL
- → `increment` — atomic counter
- → `multiply` / `divide` / `set` — atomic too
- → `push` — append to `String[]`
- → `updateMany` — bulk, `{ count }`
- → `upsert` — where + update + create

UPDATE METHODS

METHOD	NO MATCH	RETURNS
update	P2025	record
updateMany	`count: 0`	`{ count }`
updateManyAndReturn	[]	record[]
upsert	creates	record

ATOMIC OPERATIONS

OPERATION	APPLIES TO
`{ increment: n }`	Int, BigInt, Float, Decimal
`{ decrement: n }`	numbers
`{ multiply: n }`	numbers
`{ divide: n }`	numbers
`{ set: v }`	any field
`{ push: v }`	scalar lists

⚠ COMMON MISTAKES

- `views: post.views + 1` — races; use `increment`
- `update` with a non-unique `where` — use `updateMany`
- Assuming `updateMany` throws on no match — check `count`
- Upsert `where` on a non-unique field — not allowed

✓ TIPS & BEST PRACTICES

- Catch `P2025` from `update` → HTTP 404
- Optimistic lock: `where: { id, version }` + increment
- Nested: `profile: { update: { bio } }`
- Keep upsert `create` and `update` consistent

QUICK SUMMARY

- `update` — one by unique key
- `updateMany` — bulk, `{ count }`
- `increment` — race-free counter
- `upsert` — create or update

CRUD: Delete & Batch Operations

REMOVING DATA

`delete()` removes one record by unique `where`, returns it, and throws `P2025` if it's missing. `deleteMany()` removes every match in one SQL statement and returns `{ count }`. Related rows follow the relation's `onDelete` action: `Cascade` removes them, `Restrict` blocks the delete with `P2003`, `SetNull` detaches them.

NOTES

- `deleteMany()` with no `where` empties the table
- Cascades run in the DB, not in Prisma
- Deleted record is returned — log it first
- Soft delete = `update` a `deletedAt` field

DELETE EXAMPLES

```ts
await prisma.post.delete({ where: { id: 7 } });      // deleted Post, or P2025

await prisma.post.deleteMany({
  where: { published: false, createdAt: { lt: new Date("2025-01-01") } },
});                                                  // { count: 12 }

await prisma.user.delete({ where: { id: 1 } });      // Profile cascades

// Restrict FK → P2003: remove children first, atomically
await prisma.$transaction([
  prisma.post.deleteMany({ where: { authorId: 2 } }),
  prisma.user.delete({ where: { id: 2 } }),
]);

// soft delete instead of removing
await prisma.user.update({ where: { id: 3 }, data: { deletedAt: new Date() } });
await prisma.user.findMany({ where: { deletedAt: null } });   // active only

await prisma.auditLog.deleteMany();                  // ⚠ every row
```

CODE EXPLAINED

- → `delete` — one row by unique key
- → `deleteMany` + filters — one SQL DELETE
- → `onDelete: Cascade` — DB removes children
- → `$transaction([...])` — both or neither
- → `deletedAt` — keep row, hide it
- → `deletedAt: null` — filter active rows
- → no `where` — truncates the data

DELETE METHODS

METHOD	NO MATCH	RETURNS
`delete`	P2025	deleted record
`deleteMany`	`count: 0`	`{ count }`
`nested delete`	throws	parent
`nested deleteMany`	ok	parent
`loop of delete()`	—	N queries — avoid

DELETE STRATEGIES

STRATEGY	USE WHEN
`Hard delete`	Data has no recovery/audit need
`Cascade`	Children are meaningless alone
`Restrict`	Children must be handled first
`Soft delete`	Undo, audit, or legal retention
`Archive table`	Keep history, slim hot table

⚠ COMMON MISTAKES

- `deleteMany()` with an empty `where` — wipes the table
- Deleting a parent under `Restrict` — fails with `P2003`
- Loop of `delete()` for bulk removal — use `deleteMany`
- Soft-deleting but forgetting `deletedAt: null` filters

✓ TIPS & BEST PRACTICES

- Guard bulk deletes: refuse an empty `where` in services
- Return the deleted record to support an "undo" toast
- Index `deletedAt` when most queries filter on it
- Centralize soft-delete filters with `$extends` (Sheet 21)

QUICK SUMMARY

- `delete` — one by unique key, P2025
- `deleteMany` — bulk, `{ count }`
- `onDelete` — Cascade · Restrict · SetNull
- soft delete — `deletedAt` timestamp

Filtering: Where Clauses & Operators

equals · gt/gte/lt/lte · contains · startsWith · in/notIn · not · AND/OR/NOT

BUILDING FILTERS

`where` takes an object of field conditions. Multiple keys are combined with an implicit **AND**. A plain value means `equals`; an object unlocks operators like `gt`, `contains`, or `in`. Use `OR: [...]`, `AND: [...]`, and `NOT` for boolean logic — they nest to any depth and work in every query method.

NOTES

- `undefined` value = condition ignored
- `null` value = `IS NULL`
- `mode: 'insensitive'` — Postgres & Mongo
- MySQL: case-insensitive by default collation

WHERE EXAMPLES

```ts
await prisma.user.findMany({
  where: {
    email: { endsWith: "@acme.com" },          // LIKE '%@acme.com'
    age:   { gte: 18, lt: 65 },                // 18 ≤ age < 65
    role:  { in: ["ADMIN", "EDITOR"] },        // IN (...)
    name:  { not: "Bot" },                     // <> 'Bot'
  },
});
await prisma.post.findMany({
  where: {
    published: true,                           // AND (implicit)
    OR: [
      { title: { contains: "prisma", mode: "insensitive" } },
      { body:  { contains: "prisma", mode: "insensitive" } },
    ],
    NOT: { status: "ARCHIVED" },
    slug: search ? { startsWith: search } : undefined, // optional
  },
});
```

CODE EXPLAINED

- → `endsWith` — suffix match
- → `gte` + `lt` — range on one field
- → `in` — match any value in the list
- → `not` — inequality
- → `OR: [...]` — at least one branch
- → `NOT` — exclude matches
- → `undefined` — filter skipped

FIELD OPERATORS

OPERATOR	EXAMPLE	SQL
`equals`	`{ equals: 5 }` or `5`	`= 5`
`gt / gte`	`{ gt: 5 }`	`> 5` / `>= 5`
`lt / lte`	`{ lte: 5 }`	`< 5` / `<= 5`
`in / notIn`	`{ in: [1, 2] }`	`IN (1, 2)`
`contains`	`{ contains: 'a' }`	`LIKE '%a%'`
`startsWith`	`{ startsWith: 'a' }`	`LIKE 'a%'`
`endsWith`	`{ endsWith: 'a' }`	`LIKE '%a'`

LOGICAL OPERATORS

OPERATOR	MEANING
`{ a, b }`	a AND b (implicit)
`AND: [a, b]`	Explicit AND
`OR: [a, b]`	a OR b
`NOT: a`	NOT a
`NOT: [a, b]`	NOT a AND NOT b
`OR: []`	matches nothing
`AND: []`	matches everything

⚠ COMMON MISTAKES

- Expecting `contains` to be case-insensitive on Postgres — add `mode`
- `null` vs `undefined`: `null` filters, `undefined` ignores
- `OR: []` from an empty array returns no rows — guard it
- Two conditions on the same key in one object — last wins

✓ TIPS & BEST PRACTICES

- Build optional filters with `undefined` — no `if` chains
- Use `in` instead of many `OR` equals — shorter SQL
- Leading-wildcard `contains` can't use B-tree indexes
- Type reusable filters as `Prisma.UserWhereInput`

QUICK SUMMARY

- `{ field: value }` — equals
- `gt` · `gte` · `lt` · `lte` — ranges
- `contains` · `startsWith` · `endsWith` — text
- `in` · `notIn` · `not` — sets/negation
- `AND` · `OR` · `NOT` — logic

Filtering: Advanced & Pagination

ADVANCED FILTERS & PAGING

Offset pagination (`skip`/`take`) is simple but slows down as the offset grows. **Cursor pagination** starts after a known unique value, so every page costs the same. Relation filters (`some`, `every`, `none`, `is`) filter parents by their children; list filters (`has`, `hasSome`) query scalar arrays.

NOTES

- Cursor must be unique & in `orderBy`
- `skip: 1` skips the cursor row itself
- `every` is true when there are no children
- Null check: `{ field: null }` / `{ not: null }`

CURSOR PAGING + ADVANCED FILTERS

```ts
async function page(cursor?: number) {
  const rows = await prisma.post.findMany({
    take: 20,
    ...(cursor && { skip: 1, cursor: { id: cursor } }),
    orderBy: { id: "desc" },
  });
  return { rows, next: rows.at(-1)?.id };   // pass `next` back as cursor
}
await prisma.user.findMany({
  where: {
    posts: { some: { published: true } },   // ≥1 published post
    comments: { none: {} },                 // no comments at all
    profile: { isNot: null },               // 1-1 relation exists
    tags: { hasSome: ["ts", "orm"] },       // String[] overlap
    deletedAt: null,                        // IS NULL
  },
  orderBy: [
    { lastLogin: { sort: "desc", nulls: "last" } },
    { posts: { _count: "desc" } },
  ],
});
```

CODE EXPLAINED

- → `cursor` + `skip: 1` — start after last row
- → `take: 20` — page size (negative = backwards)
- → `some` — at least one child matches
- → `none: {}` — has zero children
- → `isNot: null` — optional 1-1 exists
- → `hasSome` — array overlaps list
- → `nulls: 'last'` — null ordering
- → `_count` — sort by relation count

RELATION FILTERS

FILTER	MATCHES PARENT WHEN
`some: {…}`	≥1 child matches
`every: {…}`	all children match (or none exist)
`none: {…}`	no child matches
`is: {…}`	to-one relation matches
`isNot: {…}`	to-one doesn't match
`is: null`	to-one is missing

SCALAR LIST FILTERS

FILTER	MEANING
`has: v`	contains v
`hasSome: [..]`	any overlap
`hasEvery: [..]`	contains all
`isEmpty: true`	empty array
`equals: [..]`	exact array

OFFSET VS CURSOR

	OFFSET	CURSOR
`Jump to N`	✔	✖
`Deep pages`	slow	fast
`Insert-safe`	✖	✔
`Page count`	✔	✖

⚠ COMMON MISTAKES

- Cursor on a non-unique column — rows skipped or repeated
- Forgetting `skip: 1` — the cursor row repeats
- `isEmpty` on a `String` — it's for scalar lists; use `equals: ""`
- `every` returning parents with no children — add `some: {}`

✓ TIPS & BEST PRACTICES

- Return `next` cursor in your API instead of page numbers
- Add `id` as the last `orderBy` for deterministic order
- Index the cursor/orderBy columns together
- Use offset only for small, admin-style tables

QUICK SUMMARY

- `cursor` + `skip: 1` + `take` — keyset paging
- `some` · `every` · `none` — to-many filters
- `is` · `isNot` — to-one filters
- `has` · `hasSome` · `hasEvery` — arrays
- `nulls: 'last'` — null sort order

Relations: Include & Select

include · select · nested where/orderBy/take · _count · relationLoadStrategy

LOADING RELATED DATA

`include` returns all scalar fields **plus** the listed relations. `select` returns **only** the listed fields and relations. Both nest, and a to-many relation accepts `where`, `orderBy`, `take`, and `skip` to limit what's loaded. `_count` returns relation counts without loading rows.

NOTES

- `select` ⊂ `include` — can't be siblings
- Inside `include`, use `select` on a relation
- `relationLoadStrategy` needs preview `relationJoins`
- Default `query` strategy: one query per level

INCLUDE & SELECT

```ts
const user = await prisma.user.findUnique({
  where: { id: 1 },
  include: {
    profile: true,                                  // 1-1
    posts: {                                        // filtered 1-n
      where: { published: true },
      orderBy: { createdAt: "desc" },
      take: 5,
      include: { comments: { take: 3 } },           // 2 levels deep
    },
    _count: { select: { posts: true, followers: true } },
  },
});
const cards = await prisma.post.findMany({
  relationLoadStrategy: "join",                     // single SQL query
  select: {
    id: true,
    title: true,
    author: { select: { name: true, avatarUrl: true } },
    tags: { select: { name: true } },
  },
});   // { id, title, author: { name, avatarUrl }, tags: { name }[] }[]
```

CODE EXPLAINED

- → `include: { profile: true }` — add relation
- → nested `where` / `orderBy` / `take` — limit children
- → `include` inside include — multi-level
- → `_count` — counts, no rows loaded
- → `select` — exact payload shape
- → `relationLoadStrategy` — `join` or `query`

INCLUDE VS SELECT

	INCLUDE	SELECT
Scalars returned	all	only listed
Relations	listed	listed
Payload size	larger	minimal
Best for	internal logic	API responses
Nest the other	✔	✔

NESTED RELATION ARGS

ARG	EFFECT
`where`	Filter children
`orderBy`	Sort children
`take / skip`	Limit children
`cursor`	Page children
`select / include / omit`	Shape children
`distinct`	Unique children

⚠ COMMON MISTAKES

- `include: { posts: true }` on large relations — loads every row
- `select` and `include` side by side — validation error
- Deep includes in list endpoints — huge payloads
- Reading `user.posts` without including it

✓ TIPS & BEST PRACTICES

- Always `take` on to-many includes in list views
- Use `_count` for badges instead of loading rows
- Try `relationLoadStrategy: 'join'` on Postgres for deep trees
- Share shapes with `satisfies Prisma.PostSelect` (Sheet 23)

QUICK SUMMARY

- `include` — all scalars + relations
- `select` — only what you list
- nested `where/orderBy/take` — trim children
- `_count` — relation counts

Relations: Nested Writes

WRITING THROUGH RELATIONS

Nested writes change related records inside a parent `create` or `update`. Prisma wraps the whole operation in a **transaction**, so the parent and all nested changes succeed or fail together. `connect` links an existing record by unique key; `create` makes a new one; `connectOrCreate` does whichever applies.

NOTES

- `connect` throws `P2025` if the target is missing
- `disconnect` sets the FK to NULL (optional FK)
- `set` replaces all links (to-many)
- `createMany` nested: 1-n only, no deeper nesting

NESTED WRITE EXAMPLES

```
await prisma.post.create({
  data: {
    title: "Hello",
    author: { connect: { email: "ada@example.com" } },        // existing user
    category: { connectOrCreate: {
      where: { slug: "orm" }, create: { slug: "orm", name: "ORM" } } },
  },
});
await prisma.user.update({
  where: { id: 1 },
  data: {
    profile: { upsert: { create: { bio: "Hi" }, update: { bio: "Hi" } } },
    posts: {
      update: { where: { id: 7 }, data: { published: true } },
      updateMany: { where: { published: false }, data: { views: 0 } },
      deleteMany: { createdAt: { lt: new Date("2024-01-01") } },
      disconnect: [{ id: 9 }],                                  // authorId → NULL
    },
  },
});
```

CODE EXPLAINED

- → `connect` — link by unique key
- → `connectOrCreate` — link or create
- → `upsert` — 1-1 create-or-update
- → `update` — change one child
- → `updateMany` / `deleteMany` — bulk on children
- → `disconnect` — unlink, keep row

NESTED OPERATIONS

OP	TO-ONE	TO-MANY	IN CREATE
`create`	✔	✔	✔
`createMany`	✖	✔	✔
`connect · connectOrCreate`	✔	✔	✔
`disconnect`	✔	✔	✖
`set`	✖	✔	✖
`update · upsert · delete`	✔	✔	✖
`updateMany / deleteMany`	✖	✔	✖

TWO WAYS TO LINK

STYLE	EXAMPLE
`Checked (relation)`	`author: { connect: { id: 1 } }`
`Unchecked (FK)`	`authorId: 1`
`Disconnect to-one`	`author: { disconnect: true }`
`Delete to-one`	`profile: { delete: true }`
`Replace to-many`	`tags: { set: [{ id: 2 }] }`

⚠ COMMON MISTAKES

- Mixing `authorId` and `author: {…}` in one `data`
- `disconnect` on a required FK — fails
- Concurrent `connectOrCreate` — can race to `P2002`
- `set` when you meant `connect` — drops links

✓ TIPS & BEST PRACTICES

- Prefer nested writes to manual transactions
- `connect` by natural key (`email` , `slug`)
- Retry `connectOrCreate` on `P2002`
- `include` the result to confirm changes

QUICK SUMMARY

- `create` / `connect` — new or existing child
- `connectOrCreate` — whichever applies
- `disconnect` / `set` — unlink / replace
- `update` / `upsert` / `delete` — modify children
- nested writes — always transactional

Relations: Many-to-Many & Advanced

MANY-TO-MANY

Implicit m-n uses list fields on both models; Prisma creates and manages a hidden `_AToB` join table. Use it when the link carries no data. **Explicit** m-n defines a join model with two relations and a composite key, so you can store extra columns like `grade` or `enrolledAt`.

NOTES

- Implicit join table: `_PostToTag` (A, B)
- Implicit needs single-field `@id` s on both models
- Explicit: query the join model directly
- Switching implicit → explicit needs a data migration

SCHEMA

```
// implicit: Post.tags Tag[] ↔ Tag.posts Post[] (Sheet 04)
// explicit, with data
model Student {
  id          Int          @id
  enrollments Enrollment[]
}
model Course {
  id          Int          @id
  enrollments Enrollment[]
}
model Enrollment {
  student    Student  @relation(fields: [studentId], references: [id])
  studentId  Int
  course     Course   @relation(fields: [courseId], references: [id])
  courseId   Int
  grade      String?
  enrolledAt DateTime @default(now())
  @@id([studentId, courseId])
}
```

QUERIES

```
// implicit
await prisma.post.update({
  where: { id: 1 },
  data: { tags: { connect: [{ id: 2 }, { id: 3 }] } },
});
await prisma.post.update({        // replace all
  where: { id: 1 },
  data: { tags: { set: [{ id: 4 }] } },
});
// explicit
await prisma.enrollment.create({
  data: { studentId: 1, courseId: 10, grade: "A" },
});
await prisma.student.findUnique({
  where: { id: 1 },
  include: { enrollments: { include: { course: true } } },
});
await prisma.course.findMany({    // has an A
  where: { enrollments: { some: { grade: "A" } } },
});
```

IMPLICIT VS EXPLICIT

	IMPLICIT	EXPLICIT
Join table	hidden _AToB	your model
Extra columns	✖	✔
connect / set	on either side	via join model
Query links	`tags: { some }`	`enrollments: { some }`
Uniqueness	automatic	`@@id` / `@@unique`

IMPLICIT LINK OPS

OP	EFFECT
`connect: [{id}]`	Add links
`disconnect: [{id}]`	Remove links
`set: [{id}]`	Replace all links
`set: []`	Remove all links
`create: [{…}]`	New rows + link

⚠ COMMON MISTAKES

- Adding a column to an implicit join table — use an explicit model
- No `@@id` / `@@unique` on the join model — duplicate links
- `set` when you meant `connect` — drops existing links
- Deep `include` through the join model in list views

✓ TIPS & BEST PRACTICES

- Start explicit if you might ever need link metadata
- Flatten join results: `s.enrollments.map(e => e.course)`
- `connect` to a missing ID throws — validate IDs first
- Index `courseId` on the join model for reverse lookups

QUICK SUMMARY

- implicit — `A[]` ↔ `B[]`, no extra data
- explicit — join model + `@@id`
- `connect` · `disconnect` · `set` — manage links
- `some` — filter through the link

Aggregations & Grouping

STATISTICS IN THE DATABASE

`aggregate()` computes `_count`, `_sum`, `_avg`, `_min`, and `_max` over the rows matching `where` and returns one object. `groupBy()` returns one result per distinct combination of `by` fields; `having` filters groups after aggregation. Both run as a single SQL query — no rows are loaded into Node.

NOTES

- `_sum` / `_avg` of no rows → `null`
- Decimal fields return `Decimal` objects
- `groupBy` + `skip` / `take` requires `orderBy`
- `count({ select })` counts non-null fields

AGGREGATE & GROUP

```
const stats = await prisma.order.aggregate({
  where: { status: "PAID", createdAt: { gte: new Date("2026-01-01") } },
  _count: { _all: true },
  _sum: { total: true },
  _avg: { total: true },
  _min: { total: true },
  _max: { createdAt: true },
});  // { _count: { _all: 120 }, _sum: { total: 9840.5 }, … }

const byCustomer = await prisma.order.groupBy({
  by: ["customerId", "status"],
  where: { createdAt: { gte: new Date("2026-01-01") } },
  _sum: { total: true },
  _count: { _all: true },
  having: { total: { _sum: { gt: 1000 } } },   // big spenders only
  orderBy: { _sum: { total: "desc" } },
  take: 10,
});  // [{ customerId: 4, status: "PAID", _sum: { total: 5210 }, … }]

await prisma.user.count({ select: { _all: true, phone: true } });
                                    // { _all: 50, phone: 32 }
```

CODE EXPLAINED

- → `where` — filter before aggregating
- → `_count: { _all: true }` — row count
- → `_sum` / `_avg` / `_min` / `_max` — per field
- → `by` — group key(s)
- → `having` — filter on aggregates
- → `orderBy: { _sum }` — rank groups
- → `count({ select })` — non-null counts

AGGREGATE FUNCTIONS

KEY	TYPES	NULL ROWS
`_count`	any field / _all	excluded
`_sum`	numeric	ignored
`_avg`	numeric	ignored
`_min / _max`	numeric, string, date	ignored

PRISMA → SQL

PRISMA	SQL
`aggregate({ _sum })`	`SELECT SUM(x)`
`groupBy({ by })`	`GROUP BY a, b`
`having`	`HAVING SUM(x) > n`
`distinct: ['a']`	deduped by Prisma*
`count({ where })`	`SELECT COUNT(*)`
`* nativeDistinct`	`DISTINCT ON` (PG preview)

⚠ COMMON MISTAKES

- Loading rows with `findMany` to sum in JS — use `aggregate`
- Treating `_sum` as `0` when empty — it's `null`; use `?? 0`
- Putting aggregate filters in `where` — they belong in `having`
- `groupBy` with `take` but no `orderBy` — validation error

✓ TIPS & BEST PRACTICES

- Convert `Decimal` with `.toNumber()` only for display
- Count relations with `include: { _count }` instead
- For date buckets (per day/month) use `$queryRaw` + `date_trunc`
- Index the `by` and `where` columns of hot reports

QUICK SUMMARY

- `aggregate` — one stats object
- `groupBy` — stats per group
- `having` — filter groups
- `_count` · `_sum` · `_avg` · `_min` · `_max`

Transactions

ALL-OR-NOTHING WORK

A transaction makes several queries **atomic**: if any one fails, all are rolled back. Prisma offers the **array API** — a list of queries run sequentially in one transaction — and the **interactive API**, a callback that receives a transaction client `tx` so later queries can depend on earlier results. Throwing inside the callback rolls everything back.

NOTES

- Nested writes are already transactional
- `Promise.all` is **not** a transaction
- Inside the callback use `tx`, never `prisma`
- Defaults: `maxWait` 2s · `timeout` 5s

ARRAY & INTERACTIVE TRANSACTIONS

```ts
// array API — independent queries, one transaction
const [post, count] = await prisma.$transaction([
  prisma.post.create({ data: { title: "Hi", authorId: 1 } }),
  prisma.post.count({ where: { authorId: 1 } }),
]);

// interactive API — money transfer
async function transfer(from: number, to: number, amount: number) {
  return prisma.$transaction(async (tx) => {
    const sender = await tx.account.update({
      where: { id: from },
      data: { balance: { decrement: amount } },
    });
    if (sender.balance.lt(0)) throw new Error("Insufficient funds"); // rollback
    return tx.account.update({
      where: { id: to },
      data: { balance: { increment: amount } },
    });
  }, { isolationLevel: "Serializable", maxWait: 5000, timeout: 10000 });
}
```

CODE EXPLAINED

- → array — results returned in order
- → `tx` — client bound to the transaction
- → `decrement` first — row is locked
- → `.lt(0)` — Decimal comparison
- → `throw` — rolls back both updates
- → `isolationLevel` — e.g. `Serializable`
- → `maxWait` — wait for a connection
- → `timeout` — max time the tx may run

WHICH API?

NEED	USE
Parent + children	nested write
Independent writes together	`$transaction([...])`
Read → decide → write	`$transaction(async tx => …)`
Bulk same change	`updateMany`
Long-running job	split into batches

ISOLATION LEVELS

LEVEL	PREVENTS
ReadUncommitted	nothing (dirty reads possible)
ReadCommitted	dirty reads (PG default)
RepeatableRead	+ non-repeatable reads (MySQL default)
Serializable	+ phantoms; may throw P2034

⚠ COMMON MISTAKES

- Using `prisma.` inside the callback — runs outside the transaction
- Calling HTTP APIs inside an interactive tx — holds locks and connections
- Exceeding the 5s default `timeout` — tx is rolled back with an error
- Thinking `Promise.all([...])` is atomic — it isn't

✓ TIPS & BEST PRACTICES

- Keep transactions short: prepare data before, side effects after
- Retry on `P2034` (write conflict/deadlock) with backoff
- Use atomic `increment` / `decrement` to avoid read-then-write races
- Set defaults via `transactionOptions` on the client

QUICK SUMMARY

- `$transaction([...])` — atomic batch
- `$transaction(async tx => …)` — interactive
- `throw` — rollback
- `isolationLevel` · `maxWait` · `timeout` — options

Migrations

VERSIONED SCHEMA CHANGES

Prisma Migrate compares `schema.prisma` with the database and writes SQL migration files to `prisma/migrations/`. Commit them: they are the history every environment replays. Use `migrate dev` while developing, `migrate deploy` in CI/production, and `db push` only for quick prototypes that need no history.

V7 CHANGES

- `migrate dev` no longer runs `generate`
- No auto-seed — run `prisma db seed`
- `--skip-generate` / `--skip-seed` removed
- URL comes from `prisma.config.ts`

WORKFLOW

```
# development
npx prisma migrate dev --name add_user_role    # create + apply + record
npx prisma generate                            # refresh client types
npx prisma migrate dev --create-only --name x  # write SQL, edit, then apply
npx prisma db push                             # prototype: sync, no files
npx prisma db seed                             # run migrations.seed

# CI / production
npx prisma migrate status                      # pending? drift?
npx prisma migrate deploy                      # apply pending only

# fixing history
npx prisma migrate resolve --applied 20261003120000_add_user_role
npx prisma migrate resolve --rolled-back 20261003120000_add_user_role
npx prisma migrate reset                       # ⚠ dev only: drop + reapply
```

CODE EXPLAINED

- → `--name` — folder suffix for the migration
- → `--create-only` — review/edit SQL first
- → `db push` — no migration file or history
- → `migrate status` — compare files vs DB
- → `migrate deploy` — never prompts or resets
- → `resolve` — mark a failed migration
- → `reset` — deletes all data

DEV VS PRODUCTION

COMMAND	WHERE	DESTRUCTIVE
`migrate dev`	local	may reset on drift
`db push`	prototype	may drop data
`migrate reset`	local/test	✔ drops all
`migrate deploy`	CI/prod	✖
`migrate status`	anywhere	✖
`migrate resolve`	prod fixes	✖
`db pull`	introspect	overwrites schema
`migrate diff`	anywhere	✖ (prints SQL)

PRISMA/MIGRATIONS/

PATH	PURPOSE
`2026…_init/migration.sql`	Generated SQL
`migration_lock.toml`	Locks the provider
`_prisma_migrations (table)`	Applied history + checksums
`prisma.config.ts → seed`	`tsx prisma/seed.ts`
`shadow database`	Temp DB `migrate dev` uses to detect drift
`baseline`	`--create-only` + `resolve --applied` on existing DB

⚠ COMMON MISTAKES

- Running `migrate dev` against production — it can reset
- Editing an applied `migration.sql` — checksum mismatch
- Using `db push` on a shared DB — no history to replay
- Expecting `migrate dev` to regenerate the client in v7

✓ TIPS & BEST PRACTICES

- Run `migrate deploy` in CI before the new app starts
- Use `--create-only` to add data backfills to the SQL
- Commit `migration_lock.toml` with the migrations
- Gate deploys on `migrate status` showing no drift

QUICK SUMMARY

- `migrate dev` — create + apply (dev)
- `migrate deploy` — apply pending (prod)
- `db push` — sync without history
- `migrate resolve` / `reset` — repair / wipe

Error Handling

HANDLING QUERY ERRORS

Prisma throws typed error classes from the `Prisma` namespace. Database-level failures arrive as `PrismaClientKnownRequestError` with a stable `code` (e.g. `P2002`) and a `meta` object. Check the class with `instanceof`, branch on `code`, and translate it into a domain error or HTTP status at one boundary.

NOTES

- Import `Prisma` from your generated client
- `meta` shape can vary by driver adapter
- `…OrThrow` → `P2025` when not found
- Validation errors = bug in your query shape

CATCHING & MAPPING ERRORS

```ts
import { Prisma } from "../generated/prisma/client";

export async function createUser(email: string) {
  try {
    return await prisma.user.create({ data: { email } });
  } catch (e) {
    if (e instanceof Prisma.PrismaClientKnownRequestError) {
      switch (e.code) {
        case "P2002": throw new ConflictError(`Duplicate: ${e.meta?.target}`);
        case "P2003": throw new BadRequestError("Related record missing");
        case "P2025": throw new NotFoundError();
      }
    }
    if (e instanceof Prisma.PrismaClientValidationError) {
      throw new BadRequestError("Invalid query input");
    }
    throw e;                                          // unknown → 500
  }
}
```

CODE EXPLAINED

- → `Prisma` namespace — error classes
- → `instanceof` — narrow the type
- → `e.code` — stable error code
- → `P2002` — unique constraint
- → `P2003` — foreign key failed
- → `P2025` — record not found
- → Validation — wrong args/types
- → rethrow unknown errors

COMMON ERROR CODES

CODE	MEANING	HTTP
P1001	Can't reach database	503
P2000	Value too long for column	400
P2002	Unique constraint failed	409
P2003	Foreign key constraint failed	400/409
P2025	Record not found	404
P2034	Tx write conflict / deadlock	retry

ERROR CLASSES

CLASS	WHEN
KnownRequestError	DB error with a code
UnknownRequestError	DB error, no code
ValidationError	Bad query shape/types
InitializationError	Startup/connection fail

⚠ COMMON MISTAKES

- `e.code` without `instanceof` — `e` is `unknown`
- Leaking raw Prisma messages to clients — exposes schema
- Swallowing errors with an empty `catch`
- Retrying `P2002` — it will keep failing

✓ TIPS & BEST PRACTICES

- Map errors once in middleware, not in every handler
- Retry only transient codes: `P2034`, `P1001`
- Validate input (Zod) before querying — Sheet 29
- Log `code`, `meta`, and the model with each error

QUICK SUMMARY

- `KnownRequestError` + `code` — branch
- `P2002` 409 · `P2025` 404 · `P2003` 400
- `…OrThrow` — explicit not-found
- retry transient errors only

Performance: N+1 & Indexing

THE N+1 PROBLEM

N+1 happens when you run **1 query** for a list and then **1 more per item** to load related data — 101 queries for 100 users. Fix it by loading relations in the same call (`include`/`select`), by batching IDs with `in`, or by relying on Prisma's dataloader, which batches `findUnique` calls made in the same tick.

NOTES

- Spot N+1 with `log: ['query']`
- `findUnique` in the same tick → 1 query
- `join` strategy: preview `relationJoins`
- Indexes fix slow queries, not query count

N+1 AND FIXES

```
// ✖ N+1: 1 + N queries
const users = await prisma.user.findMany();
for (const u of users) {
  u.posts = await prisma.post.findMany({ where: { authorId: u.id } });
}

// ✔ include: relations loaded with the list
await prisma.user.findMany({ include: { posts: true } });

// ✔ manual batch: 2 queries total
const ids = users.map((u) => u.id);
const posts = await prisma.post.findMany({ where: { authorId: { in: ids } } });

// ✔ dataloader: same-tick findUnique calls are batched (GraphQL resolvers)
await Promise.all(ids.map((id) => prisma.user.findUnique({ where: { id } }).posts()));

// ✔ single SQL query with LATERAL JOINs (Postgres/MySQL)
await prisma.user.findMany({ relationLoadStrategy: "join", include: { posts: true } });
```

CODE EXPLAINED

- → loop query — one round-trip per user
- → `include` — Prisma loads posts in bulk
- → `in: ids` — one query for all children
- → fluent `.posts()` — batched by dataloader
- → `relationLoadStrategy: 'join'` — 1 query
- → default `query` strategy — 1 query/level

SYMPTOM → FIX

SYMPTOM	FIX
Query count grows with rows	`include` / `in` batch
Slow single query	`@@index` on where/orderBy
Huge payloads	`select` fewer fields
Slow deep pages	cursor pagination
Slow counts on big tables	cache or approximate
Sum in JS	`aggregate` / `groupBy`

INDEX TRADE-OFFS

FACTOR	EFFECT
Reads (WHERE/ORDER/JOIN)	faster
Writes (INSERT/UPDATE)	slower per index
Storage	grows per index
Low-cardinality column	little benefit
Composite order	leftmost prefix

⚠ COMMON MISTAKES

- `await` inside `for…of` over query results — classic N+1
- `include` everything "just in case" — slow, heavy
- Adding indexes without checking the query plan
- Fixing N+1 with caching instead of batching

✓ TIPS & BEST PRACTICES

- Log query counts per request in development
- Use `EXPLAIN ANALYZE` via `$queryRaw` on slow queries
- Benchmark `join` vs `query` strategy on your data
- Index FK columns used by `include` and filters

QUICK SUMMARY

- N+1 — 1 list query + N child queries
- `include` / `select` — load relations together
- `in: ids` — manual batch
- `@@index` — faster filters and sorts

Connection Pooling & Optimization

MANAGING CONNECTIONS

In Prisma 7 the **driver adapter owns the pool**. Pool size and timeouts are adapter options (`max`, `connectionTimeoutMillis`, `idleTimeoutMillis` for `pg`); the old `connection_limit` URL parameters no longer apply. Every app instance has its own pool, so total connections = instances × `max`, which must stay below the database's `max_connections`.

NOTES

- `pg` default `max` is 10 (v6: CPUs × 2 + 1)
- Serverless: small `max` + external pooler
- CLI uses the URL in `prisma.config.ts`
- App uses the adapter's `connectionString`

POOL CONFIGURATION & MONITORING

```ts
import { PrismaPg } from "@prisma/adapter-pg";

const adapter = new PrismaPg({
  connectionString: process.env.DATABASE_URL,   // pooled URL (PgBouncer)
  max: 10,                                      // pool size per instance
  connectionTimeoutMillis: 5_000,               // fail fast when pool is full
  idleTimeoutMillis: 30_000,                    // close idle connections
});

export const prisma = new PrismaClient({
  adapter,
  log: [{ emit: "event", level: "query" }, "warn", "error"],
});

prisma.$on("query", (e) => {
  if (e.duration > 200) console.warn(`slow ${e.duration}ms: ${e.query}`);
});

// prisma.config.ts → datasource: { url: env("DIRECT_URL") }  // migrations
```

CODE EXPLAINED

- → `max` — connections per process
- → `connectionTimeoutMillis` — default 0 = wait forever
- → `idleTimeoutMillis` — default 10s
- → `emit: 'event'` — enables `$on('query')`
- → `e.duration` — ms per query
- → `DIRECT_URL` — migrations bypass the pooler

V6 URL PARAM → V7 PG OPTION

V6	V7 (ADAPTER-PG)
`connection_limit`	`max` (default 10)
`connect_timeout`	`connectionTimeoutMillis` (0)
`max_idle_connection_lifetime`	`idleTimeoutMillis` (10s)

SIZING GUIDE

DEPLOYMENT	POOL
Single server	max 10–20
N replicas	max ≤ DB limit ÷ N
Serverless functions	max 1–5 + pooler
Edge runtime	HTTP driver / Accelerate
Monitor (PG)	pg_stat_activity

⚠ COMMON MISTAKES

- Copying v6 `?connection_limit=` into v7 URLs — ignored
- Large `max` × many replicas — "too many connections"
- Running migrations through a transaction-mode pooler
- Leaving `connectionTimeoutMillis` at 0 — requests hang

✓ TIPS & BEST PRACTICES

- Set a finite `connectionTimeoutMillis` to surface pool exhaustion
- Use a pooled URL in the app and a direct URL for the CLI
- Log slow queries with `$on('query')` and `duration`
- Call `$disconnect()` on shutdown to drain the pool

QUICK SUMMARY

- adapter options — pool size & timeouts
- instances × `max` < `max_connections`
- pooler for serverless, direct URL for CLI
- `$on('query')` — slow-query logging

Advanced: Soft Deletes & Auditing

RECOVERABLE DELETES & CHANGE HISTORY

A **soft delete** marks a row with a `deletedAt` timestamp instead of removing it, so it can be restored and audited. An **audit log** records who changed what and when. A query extension (`$extends`) can apply both automatically so individual queries don't have to remember.

NOTES

- Unique columns still collide with soft-deleted rows
- Cascades don't run on soft delete
- Extensions don't cover `$queryRaw`
- Audit inside the same transaction

SOFT-DELETE EXTENSION + AUDIT WRITE

```
export const db = prisma.$extends({
  query: {
    user: {
      async findMany({ args, query }) {
        args.where = { deletedAt: null, ...args.where };   // hide deleted
        return query(args);
      },
      async delete({ args }) {                             // delete → update
        return prisma.user.update({ where: args.where, data: { deletedAt: new Date() } });
      },
    },
  },
});
await db.user.delete({ where: { id: 5 } });                // soft delete
await prisma.user.update({ where: { id: 5 }, data: { deletedAt: null } }); // restore
await prisma.$transaction(async (tx) => {                  // change + audit
  const after = await tx.post.update({ where: { id: 9 }, data: { title: "New" } });
  await tx.auditLog.create({
    data: { model: "Post", recordId: 9, action: "UPDATE", actorId: me.id, after } });
});
```

CODE EXPLAINED

- → `$extends({ query })` — intercept queries
- → spread `args.where` — caller can override
- → `delete` → `update` — sets `deletedAt`
- → `deletedAt: null` — restore the row
- → `$transaction` — change + log together
- → `after` (and `before`) — `Json` snapshots
- → `actorId` — who made the change

SCHEMA

```
model User {
  id        Int       @id @default(autoincrement())
  deletedAt DateTime?           // null = active
  updatedAt DateTime  @updatedAt
  @@index([deletedAt])
}
model AuditLog {
  id       Int      @id @default(autoincrement())
  model    String   //   "Post"
  recordId Int
  action   String          // CREATE | UPDATE | DELETE
  actorId  Int?                 // who
  before   Json?
  after    Json?
  at       DateTime @default(now())
}
```

SOFT-DELETE COVERAGE

METHOD	REWRITE TO
`findMany / findFirst`	add `deletedAt: null`
`count / aggregate`	add `deletedAt: null`
`findUnique`	filter after or use findFirst
`delete`	`update` set `deletedAt`
`deleteMany`	`updateMany` set `deletedAt`
`purge (GDPR)`	real `delete` on base client

⚠ COMMON MISTAKES

- `@unique` email blocks re-signup after soft delete
- Filtering only `findMany` — `count` still sees deleted rows
- Writing audit rows outside the transaction — logs can lie

✓ TIPS & BEST PRACTICES

- Use a partial unique index (raw SQL migration) for active rows
- Keep a base client for admin/purge paths
- Index `deletedAt` or put it first in composite indexes
- Strip secrets before saving audit snapshots

QUICK SUMMARY

- `deletedAt` — soft delete marker
- `$extends` query — automatic filtering
- restore — set `deletedAt: null`
- `AuditLog` — model, recordId, action, before/after

Advanced: $extends & Custom Methods

CLIENT EXTENSIONS

`$extends` returns a **new client** with extra behavior and full type safety; the original stays unchanged. Four components: `model` adds methods to delegates, `query` intercepts operations (before/after), `result` adds computed fields, and `client` adds top-level methods. Extensions replace the `$use` middleware, which was removed in Prisma 7.

NOTES

- Extensions are pure TypeScript — no SQL changes
- Chain: `prisma.$extends(a).$extends(b)`
- `$allModels` / `$allOperations` wildcards
- Export the extended client as your `db`

ALL FOUR COMPONENTS

```ts
export const db = prisma.$extends({
  model: { user: {
    signUp: (email: string) => prisma.user.create({ data: { email } }),
  } },
  result: { user: {
    fullName: {                                  // computed field
      needs: { first: true, last: true },
      compute: (u) => `${u.first} ${u.last}`,
    },
  } },
  query: { $allModels: {
    async $allOperations({ model, operation, args, query }) {
      const t = performance.now();
      const res = await query(args);             // before → run → after
      console.log(`${model}.${operation}`, performance.now() - t);
      return res;
    },
  } },
  client: { $health: () => prisma.$queryRaw`SELECT 1` },
});

const u = await db.user.signUp("ada@example.com");
u.fullName;                                      // "Ada Lovelace" (typed)
```

CODE EXPLAINED

- → `model` — custom delegate methods
- → `result` — virtual field, typed
- → `needs` — fields required to compute
- → `$allModels` — apply to every model
- → `$allOperations` — every method
- → `query(args)` — run the real query
- → `client` — top-level helper

COMPONENTS

COMPONENT	USE FOR
`model`	Repository methods: `db.user.signUp()`
`query`	Logging, soft delete, tenant scoping
`result`	Computed fields, masking
`client`	Helpers: `db.$health()`

HELPERS

API	PURPOSE
`Prisma.defineExtension`	Reusable, shareable extension
`Prisma.getExtensionContext(this)`	Current model in `$allModels`
`Prisma.Exact<A, B>`	Type-safe generic args
`$use (removed v7)`	Migrate to `query` extension

⚠ COMMON MISTAKES

- Calling `db` inside its own `query` hook — infinite recursion
- Using the base `prisma` and expecting extension behavior
- Heavy `result.compute` on list queries — runs per row
- Still calling `prisma.$use()` after upgrading to v7

✓ TIPS & BEST PRACTICES

- Package extensions with `Prisma.defineExtension` in their own files
- Use `result` with `needs` to mask secrets (e.g. last 4 digits)
- Put tenant scoping in a `query` extension
- Type the extended client: `type DB = typeof db`

QUICK SUMMARY

- `$extends` — new, typed client
- `model` · `query` · `result` · `client`
- `$allModels` / `$allOperations` — wildcards
- replaces removed `$use` middleware

TypeScript & Type Safety

GENERATED TYPES

`prisma generate` emits a type for every model (`User`), every input (`Prisma.UserCreateInput`), and helpers to derive result types from a query shape. Query results are inferred automatically; when you need to name a type — for a function parameter or API response — derive it from the same `select`/`include` object so it never drifts.

NOTES

- `User` = scalars only, no relations
- `satisfies` keeps literal types for reuse
- `Prisma.validator` still works (older style)
- `…GetPayload<T>` — not "PayloadOf"

DERIVING TYPES FROM QUERIES

```ts
import type { User, Prisma } from "../generated/prisma/client";

// reusable query shape + its exact result type
const userWithPosts = {
  include: { posts: { select: { id: true, title: true } } },
} satisfies Prisma.UserDefaultArgs;

type UserWithPosts = Prisma.UserGetPayload<typeof userWithPosts>;
// User & { posts: { id: number; title: string }[] }

export function findAuthor(id: number): Promise<UserWithPosts | null> {
  return prisma.user.findUnique({ where: { id }, ...userWithPosts });
}

// typed inputs for service functions
function search(where: Prisma.PostWhereInput, orderBy?: Prisma.PostOrderByWithRelationInput) {
  return prisma.post.findMany({ where, orderBy });
}

// result type of any delegate call
type Users = Prisma.Result<typeof prisma.user, { select: { email: true } }, "findMany">;
```

CODE EXPLAINED

- → `import type` — model + namespace types
- → `satisfies` — checked, still literal
- → `UserGetPayload` — type from the shape
- → spread the shape — query & type match
- → `PostWhereInput` — typed filters
- → `OrderByWithRelationInput` — typed sort
- → `Prisma.Result` — any operation's result

GENERATED TYPE FAMILIES

TYPE	USE
`User`	Model scalars
`Prisma.UserCreateInput`	`create` data (relations)
`Prisma.UserUncheckedCreateInput`	`create` data (FK ids)
`Prisma.UserUpdateInput`	`update` data
`Prisma.UserWhereInput`	Filters
`Prisma.UserWhereUniqueInput`	Unique lookups
`Prisma.UserSelect / Include`	Shapes

TYPE HELPERS

HELPER	GIVES
`UserGetPayload<A>`	Result for args A
`Prisma.Result<D, A, Op>`	Result of an operation
`Prisma.Args<D, Op>`	Args of an operation
`Prisma.validator<T>()`	Typed shape (legacy)
`Role (enum export)`	Enum union
`Prisma.JsonValue`	Json field type

⚠ COMMON MISTAKES

- Typing results as `User` when you `include` relations
- Hand-writing interfaces that drift from the schema
- Annotating a shape (`: Prisma.UserSelect`) — widens to booleans
- Using `any` for `where` — loses filter checking

✓ TIPS & BEST PRACTICES

- Co-locate shape + `GetPayload` type in one module
- Use `satisfies`, not type annotations, on shapes
- Use `Unchecked*Input` when you have raw FK IDs
- Run `tsc --noEmit` in CI after `prisma generate`

QUICK SUMMARY

- generated types — models, inputs, enums
- `satisfies Prisma.*Args` — reusable shape
- `*GetPayload<typeof shape>` — result type
- `*WhereInput` / `*CreateInput` — inputs

Testing & Mocking

UNIT VS INTEGRATION TESTS

Unit tests replace the Prisma client with a deep mock (`vitest-mock-extended` or `jest-mock-extended`) so business logic runs without a database. **Integration tests** use a real, separate test database, apply migrations before the run, and clean tables between tests to stay isolated.

NOTES

- Never point tests at your dev DB
- Mocks can't catch SQL or constraint bugs
- Run test files serially against one DB
- Or one DB/schema per worker

MOCKED UNIT TEST (VITEST)

```ts
// src/lib/__mocks__/prisma.ts — used automatically by vi.mock()
import { beforeEach } from "vitest";
import { mockDeep, mockReset } from "vitest-mock-extended";
import type { PrismaClient } from "../../../generated/prisma/client";
export const prisma = mockDeep<PrismaClient>();
beforeEach(() => mockReset(prisma));

// src/user.service.test.ts
import { vi, it, expect } from "vitest";
import { prisma } from "./lib/__mocks__/prisma";       // typed mock
import { createUser } from "./user.service";
vi.mock("./lib/prisma");                               // swap the singleton
it("creates a user", async () => {
  prisma.user.create.mockResolvedValue({ id: 1, email: "a@b.c", name: null });
  await expect(createUser("a@b.c")).resolves.toMatchObject({ id: 1 });
  expect(prisma.user.create).toHaveBeenCalledWith({ data: { email: "a@b.c" } });
});
```

CODE EXPLAINED

- → `mockDeep` — every delegate/method mocked
- → `__mocks__/` — auto-used by `vi.mock`
- → import the mock file — typed mock API
- → `mockReset` — clean state per test
- → `mockResolvedValue` — fake DB result
- → `toHaveBeenCalledWith` — assert the query

INTEGRATION SETUP

```ts
// .env.test → DATABASE_URL=".../app_test"
// globalSetup.ts
import { execSync } from "node:child_process";
export default () =>
  execSync("npx prisma migrate reset --force", { stdio: "inherit" });

// per test: children before parents
afterEach(async () => {
  await prisma.$transaction([
    prisma.post.deleteMany(),
    prisma.user.deleteMany(),
  ]);
});

// factory
const makeUser = (o: Partial<Prisma.UserCreateInput> = {}) =>
  prisma.user.create({ data: { email: `u${Date.now()}@t.io`, ...o } });
```

ISOLATION STRATEGIES

STRATEGY	TRADE-OFF
`deleteMany per table`	simple; mind FK order
`TRUNCATE … CASCADE`	fast; via `$executeRawUnsafe`
`migrate reset`	clean; slow per test
`DB/schema per worker`	parallel; more setup
`Testcontainers`	real DB in Docker

⚠ COMMON MISTAKES

- Mock tests only — constraint and SQL bugs reach prod
- Deleting parents before children — FK error `P2003`
- Parallel test files sharing one DB — flaky results
- `migrate reset` without `--force` in CI — prompts and hangs

✓ TIPS & BEST PRACTICES

- Load `.env.test` with `dotenv` before Prisma is imported
- Use factories with unique emails to avoid `P2002`
- Keep the mock module path identical to the real import
- Run integration tests in CI against a disposable DB

QUICK SUMMARY

- `mockDeep<PrismaClient>()` — unit tests
- test DB + `migrate reset --force` — integration
- `deleteMany` / TRUNCATE — isolation
- factories — quick valid records

Debugging & Studio

SEEING WHAT PRISMA DOES

Prisma Studio is a local GUI for browsing and editing data (`npx prisma studio`, port 5555). For code, enable **query logging** to see the SQL, parameters, and duration of each query, and use the DEBUG environment variable for internal client/CLI logs. CLI commands like `validate` and `debug` catch configuration problems before runtime.

NOTES

- Studio edits real data — careful on prod
- `log: ['query']` prints to stdout
- `emit: 'event'` → handle with `$on`
- Params are logged separately from SQL

LOGGING & INSPECTION

```
const prisma = new PrismaClient({
  adapter,
  log: [
    { emit: "event", level: "query" },
    { emit: "stdout", level: "warn" },
    { emit: "stdout", level: "error" },
  ],
});

prisma.$on("query", (e) => {
  console.log(`${e.duration}ms ${e.query}`);    // SELECT … WHERE "id" = $1
  console.log("params:", e.params);              // ["42"]
});

// time + label every operation (extension)
const traced = prisma.$extends({ query: {
  async $allOperations({ model, operation, args, query }) {
    console.time(`${model}.${operation}`);
    try { return await query(args); }
    finally { console.timeEnd(`${model}.${operation}`); }
  },
} });
```

CODE EXPLAINED

- → `emit: 'event'` — route logs to `$on`
- → `emit: 'stdout'` — print directly
- → `e.query` — generated SQL
- → `e.params` — bound values
- → `e.duration` — milliseconds
- → `$allOperations` — per-op timing

DEBUG CLI

COMMAND	DOES
`npx prisma studio`	GUI at localhost:5555
`prisma studio --port 5556`	Custom port
`npx prisma validate`	Check schema
`npx prisma format`	Format + fix relations
`npx prisma debug`	Print config/env info
`npx prisma -v`	Versions

ENVIRONMENT & TOOLS

SETTING	EFFECT
`DEBUG="prisma:*"`	All internal debug logs
`DEBUG="prisma:client"`	Client logs only
`errorFormat: 'pretty'`	Readable error output
`@prisma/instrumentation`	OpenTelemetry spans
`EXPLAIN ANALYZE`	Plan via `$queryRaw`

⚠ COMMON MISTAKES

- Logging `query` events in production — noisy, may leak data
- Reading `e.query` without `e.params` — values look missing
- Running Studio against production casually
- Debugging stale types — rerun `prisma generate` first

✓ TIPS & BEST PRACTICES

- Log only queries slower than a threshold in production
- Add tracing with `@prisma/instrumentation` + OpenTelemetry
- Use Studio to verify seeds and migrations locally
- Paste slow SQL into `EXPLAIN ANALYZE`

QUICK SUMMARY

- `prisma studio` — data GUI
- `log` + `$on('query')` — SQL + timing
- `DEBUG=prisma:*` — internal logs
- `validate` · `format` · `debug` — CLI checks

Raw SQL Queries

ESCAPE HATCH TO SQL

Use raw SQL when Prisma's API can't express a query: CTEs, window functions, `date_trunc` buckets, full-text search, or vendor features. `$queryRaw` returns rows; `$executeRaw` returns the affected-row count. Written as **tagged templates**, every `${value}` becomes a bound parameter, so input is never spliced into SQL. **TypedSQL** adds generated types for `.sql` files.

NOTES

- Results are not validated — type is a cast
- Use DB column names, not `@map`ped names
- `bigint` → `BigInt`, `numeric` → `Decimal`
- TypedSQL: preview feature `typedSql`

RAW QUERY EXAMPLES

```ts
import { Prisma } from "../generated/prisma/client";

const email = "ada@example.com";
const users = await prisma.$queryRaw<{ id: number; email: string }[]>`
  SELECT id, email FROM users WHERE email = ${email}`;   // $1 param

const n = await prisma.$executeRaw`
  UPDATE posts SET views = 0 WHERE created_at < ${cutoff}`;  // 37

const ids = [1, 2, 3];
await prisma.$queryRaw`SELECT * FROM users WHERE id IN (${Prisma.join(ids)})`;

const byStatus = status ? Prisma.sql`AND status = ${status}` : Prisma.empty;
await prisma.$queryRaw`SELECT * FROM orders WHERE total > ${min} ${byStatus}`;

const daily = await prisma.$queryRaw`
  SELECT date_trunc('day', created_at) AS day, count(*)::int AS n
  FROM orders GROUP BY 1 ORDER BY 1`;

await prisma.$queryRawUnsafe("SELECT * FROM users WHERE id = $1", id); // positional
```

CODE EXPLAINED

- → `$queryRaw<T>` — rows cast to `T`
- → `${email}` — parameter, not concatenation
- → `$executeRaw` — affected row count
- → `Prisma.join` — safe `IN (...)` list
- → `Prisma.sql` / `Prisma.empty` — compose
- → `::int` — avoid `BigInt` from `count(*)`
- → `…Unsafe` — string SQL; params only

TYPEDSQL

```sql
-- prisma/sql/topAuthors.sql
-- @param {Int} $1:minPosts
SELECT u.id, u.email, count(p.id)::int AS posts
FROM users u JOIN posts p ON p.author_id = u.id
GROUP BY u.id HAVING count(p.id) >= $1;

-- npx prisma generate --sql
-- import { topAuthors } from "../generated/prisma/sql";
-- await prisma.$queryRawTyped(topAuthors(5));
```

SAFE VS UNSAFE

API	INJECTION-SAFE?
`$queryRaw`…${v}``	✔ parameterized
`Prisma.sql`…${v}``	✔ parameterized
`Prisma.join(arr)`	✔ parameterized
`$queryRawUnsafe(sql, v)`	✔ only with $1 params
`$queryRawUnsafe(`…${v}`)`	✖ string concat
`Prisma.raw(str)`	✖ inserted verbatim

⚠ COMMON MISTAKES

- `$queryRawUnsafe` with template-string interpolation — SQL injection
- Quoting params yourself: `'${v}'` — breaks binding
- Interpolating table/column names — params can't be identifiers
- Using model field names instead of mapped column names

✓ TIPS & BEST PRACTICES

- Prefer TypedSQL for complex, frequently used queries
- Cast `count(*)` to `::int` to get `number`
- Allow-list identifiers before using `Prisma.raw`
- Use `EXPLAIN ANALYZE` the same way for tuning

QUICK SUMMARY

- `$queryRaw` — SELECT → rows
- `$executeRaw` — writes → count
- tagged template — automatic parameters
- `Prisma.join` / `Prisma.sql` — safe composition
- TypedSQL — `.sql` files with types

Sheet 27

Special Cases: Null, JSON, Arrays

null · Prisma.JsonNull · DbNull · path · String[] · BigInt · Bytes · Decimal

EDGE-CASE TYPES

Optional fields hold `null`, but `Json` fields have **two** nulls: a database `NULL` (`Prisma.DbNull`) and a JSON `null` value (`Prisma.JsonNull`). JSON fields can be filtered by `path`. Scalar lists (`String[]`) are native on Postgres/CockroachDB. `BigInt` maps to JS `bigint`, `Bytes` to `Uint8Array`, and `Decimal` to a `Decimal` object.

NOTES

- `Prisma.AnyNull` filters both nulls
- Postgres `path` = array; MySQL = `'$.a.b'`
- `JSON.stringify` throws on `bigint`
- Lists: no MySQL/SQLite — use Json

NULL, JSON & LISTS

```ts
import { Prisma } from "../generated/prisma/client";

await prisma.user.findMany({ where: { phone: null } });          // IS NULL
await prisma.user.findMany({ where: { phone: { not: null } } }); // IS NOT NULL

await prisma.user.create({ data: { email: "a@b.c", settings: { theme: "dark", beta: true } } });
await prisma.user.update({ where: { id: 1 }, data: { settings: Prisma.DbNull } });   // SQL NULL
await prisma.user.update({ where: { id: 2 }, data: { settings: Prisma.JsonNull } }); // JSON null

await prisma.user.findMany({
  where: { settings: { path: ["theme"], equals: "dark" } },      // Postgres
});
await prisma.product.findMany({
  where: { meta: { path: ["tags"], array_contains: ["sale"] } },
});

await prisma.post.update({ where: { id: 1 }, data: { tags: { push: "orm" } } });
await prisma.post.findMany({ where: { tags: { hasEvery: ["orm", "ts"] } } });

const u = await prisma.user.findUnique({ where: { id: 1 } });
JSON.stringify(u, (_, v) => (typeof v === "bigint" ? v.toString() : v));
```

CODE EXPLAINED

- → `null` — IS NULL; `not: null` — IS NOT NULL
- → plain object — stored as JSON
- → `DbNull` — column NULL
- → `JsonNull` — JSON literal `null`
- → `path` + `equals` — nested key match
- → `array_contains` — JSON array check
- → `push` / `hasEvery` — scalar lists
- → replacer — serialize `BigInt`

JSON FILTERS

FILTER	MATCHES
`equals`	exact value at path
`string_contains`	substring at path
`string_starts_with`	prefix at path
`array_contains`	array includes values
`gt / lt (at path)`	numeric compare
`equals: Prisma.AnyNull`	DB or JSON null

SPECIAL TYPES

TYPE	JS VALUE / GOTCHA
`BigInt`	`10n` — not JSON-serializable
`Bytes`	`Uint8Array` (v6+, was Buffer)
`Decimal`	`new Prisma.Decimal('9.99')`
`DateTime`	`Date`, stored UTC
`String[]`	`set · push · has · isEmpty`
`Unsupported`	raw SQL only

⚠ COMMON MISTAKES

- Setting a Json field to `null` — use `DbNull` or `JsonNull`
- Using MySQL path syntax on Postgres (or vice versa)
- Returning `BigInt` from an API — `JSON.stringify` throws
- Comparing `Decimal` s with `===` — use `.equals()`

✓ TIPS & BEST PRACTICES

- Validate JSON shape with Zod on write (Sheet 29)
- Add a GIN index for frequent JSON path filters
- Prefer `Json` columns for flexible data, columns for queried data
- Convert `Decimal` with `.toFixed(2)` for display

QUICK SUMMARY

- `null` / `not: null` — nullable fields
- `DbNull` vs `JsonNull` — two Json nulls
- `path` filters — query inside JSON
- `String[]` — `push` · `has` · `hasSome`
- `BigInt` · `Bytes` · `Decimal` — special JS types

Self-Referential & Complex Relations

MODELS THAT POINT TO THEMSELVES

A self-relation links a model to itself: an employee's manager, a category's parent, or users following users. Because both ends are the same model, each pair of fields needs a shared **relation name**. Prisma loads a fixed number of levels with nested `include`; for unknown depth, use a recursive CTE with `$queryRaw`.

NOTES

- Self 1-n: FK + named `@relation` pair
- Self m-n: two list fields, one name
- SQL Server: self-relations need `NoAction`
- Guard against cycles in app logic

SCHEMA

```
model Employee {
  id        Int        @id @default(autoincrement())
  managerId Int?
  manager   Employee?  @relation("Reports", fields: [managerId], references: [id])
  reports   Employee[] @relation("Reports")
}
model Category {
  id       Int        @id @default(autoincrement())
  parentId Int?
  parent   Category?  @relation("Tree", fields: [parentId], references: [id], onDelete: Cascade)
  children Category[] @relation("Tree")
  @@index([parentId])
}
model User {
  id        Int    @id @default(autoincrement())
  followers User[] @relation("Follows")
  following User[] @relation("Follows")
}
```

CODE EXPLAINED

- → `"Reports"` — names the field pair
- → `managerId Int?` — the root has no manager
- → `reports Employee[]` — virtual list
- → `onDelete: Cascade` — deletes the subtree
- → `@@index([parentId])` — fast child lookups
- → `followers` / `following` — self m-n

QUERIES

```
await prisma.employee.findUnique({
  where: { id: 1 },
  include: { manager: true, reports: { include: { reports: true } } },
});                                        // 2 levels down
await prisma.user.update({                 // follow
  where: { id: 1 },
  data: { following: { connect: { id: 2 } } },
});
// unknown depth: recursive CTE
await prisma.$queryRaw`
  WITH RECURSIVE tree AS (
    SELECT id, "parentId", 0 AS depth FROM "Category" WHERE id = ${rootId}
    UNION ALL
    SELECT c.id, c."parentId", t.depth + 1
    FROM "Category" c JOIN tree t ON c."parentId" = t.id
  ) SELECT * FROM tree`;
```

PATTERNS & DEPTH

PATTERN	SHAPE / TRADE-OFF
Hierarchy (1-n)	manager? ↔ reports[]
Tree	parent? ↔ children[]
Follows (m-n)	followers[] ↔ following[]
Nested include	fixed depth, typed
WITH RECURSIVE	any depth, 1 query, untyped
Materialized path	'/1/4/9' + startsWith

⚠ COMMON MISTAKES

- Unnamed self-relations — ambiguity error in the schema
- Cascade on deep trees — deletes whole subtrees
- Unbounded recursion on cyclic data — add a depth limit
- Deep nested `include` — exponential payload growth

✓ TIPS & BEST PRACTICES

- Index `parentId` / `managerId` for child lookups
- Use materialized paths for read-heavy category trees
- Use an explicit `Follow` model to store `createdAt`
- Prevent cycles before setting a new parent

QUICK SUMMARY

- `@relation("Name")` — pair self-relation fields
- `parent?` / `children[]` — tree
- `followers` / `following` — self m-n
- `WITH RECURSIVE` — unlimited depth

Validation & Custom Logic

zod · safeParse · z.infer · $extends query · CHECK constraint · business rules

VALIDATING BEFORE YOU WRITE

Prisma's types check **shape** at compile time, not **content** at runtime: an empty string or a negative price still reaches the database. Validate untrusted input with a schema library such as **Zod** at the API boundary, enforce business rules in services, and keep database constraints (`@unique`, FKs, CHECK) as the final safety net.

NOTES

- Validate at the edge: request → DTO → query
- Pre-save hooks: `query` extension
- `CHECK` constraints via a SQL migration
- `zod-prisma-types` can generate schemas

ZOD + EXTENSION HOOK

```ts
import { z } from "zod";

export const ProductInput = z.object({
  name: z.string().trim().min(2).max(120),
  sku: z.string().regex(/^[A-Z0-9-]{4,20}$/),
  price: z.number().positive().multipleOf(0.01),
  stock: z.number().int().min(0).default(0),
});
export type ProductInput = z.infer<typeof ProductInput>;

export async function createProduct(body: unknown) {
  const r = ProductInput.safeParse(body);
  if (!r.success) throw new BadRequestError(z.prettifyError(r.error));
  if (await prisma.product.count({ where: { sku: r.data.sku } }))
    throw new ConflictError("SKU exists");          // business rule
  return prisma.product.create({ data: r.data });
}

export const db = prisma.$extends({ query: { product: {
  async update({ args, query }) {                   // pre-save hook
    args.data = ProductInput.partial().parse(args.data);
    return query(args);
  },
} } });
```

CODE EXPLAINED

- → `z.object` — runtime schema
- → `.trim().min()` — normalize + check
- → `z.infer` — TS type from schema
- → `safeParse` — no throw; inspect result
- → `prettifyError` — readable message (Zod 4)
- → business rule — service-level check
- → `query` extension — hook on every update

WHERE VALIDATION LIVES

LAYER	CATCHES
TypeScript types	wrong shape at compile time
Zod at API edge	bad input, formats, ranges
Service logic	business rules, permissions
Extension hook	rules for every write path
DB constraints	races, other writers

CHECK CONSTRAINT (MIGRATION.SQL)

```sql
-- npx prisma migrate dev --create-only --name price_check
ALTER TABLE "Product"
  ADD CONSTRAINT price_positive CHECK (price > 0);
```

⚠ COMMON MISTAKES

- Trusting TS types for request bodies — runtime data is `unknown`
- Count-then-create uniqueness checks — race; still catch `P2002`
- Validating only in the UI — APIs get called directly
- Hooking `update` but not `updateMany` / `upsert`

✓ TIPS & BEST PRACTICES

- Return field-level errors from `error.issues` to the client
- Reuse one Zod schema for create; `.partial()` for update
- Normalize input (trim, lowercase emails) in the schema
- Encode critical invariants as DB constraints too

QUICK SUMMARY

- `zod` — runtime validation at the edge
- `z.infer` — single source for types
- `$extends` query — pre-save hooks
- DB constraints — final safety net

Best Practices & Patterns

PRODUCTION-READY PRISMA

Keep Prisma behind a thin **data-access layer**: one client instance, query shapes defined once, errors mapped once, and services that never leak Prisma types to HTTP responses. Split large schemas into **multiple `.prisma` files** by domain, and keep naming consistent so the generated API reads naturally.

NAMING

- Model: `PascalCase` singular · field: `camelCase`
- FK: `authorId` next to `author`
- Booleans: `isActive`, `hasPaid`
- Enums: `PascalCase`; values `UPPER_CASE`

PROJECT LAYOUT

```
prisma/
  schema/                 # multi-file schema
    base.prisma           # generator + datasource
    user.prisma
    order.prisma
  migrations/
  seed.ts
prisma.config.ts          # schema: "prisma/schema"
generated/prisma/         # gitignored client
src/
  lib/prisma.ts           # singleton + $extends
  lib/prisma-errors.ts    # P-code → HTTP mapping
  modules/user/
    user.repo.ts          # queries + shapes
    user.service.ts       # rules, transactions
    user.schema.ts        # zod input schemas
```

REPOSITORY PATTERN

```ts
// user.repo.ts
export const publicUser = {
  select: { id: true, email: true, name: true, createdAt: true },
} satisfies Prisma.UserDefaultArgs;
export type PublicUser = Prisma.UserGetPayload<typeof publicUser>;

export const userRepo = {
  byId: (id: number) =>
    db.user.findUnique({ where: { id }, ...publicUser }),
  list: (cursor?: number, take = 20) =>
    db.user.findMany({
      ...publicUser, take,
      ...(cursor && { skip: 1, cursor: { id: cursor } }),
      orderBy: { id: "desc" },
    }),
  create: (data: Prisma.UserCreateInput) =>
    db.user.create({ data, ...publicUser }),
};
```

DO / AVOID

DO	AVOID
one `PrismaClient`	client per request
`select` shapes	returning full rows to APIs
cursor pagination	deep `skip` offsets
`createMany` / `updateMany`	awaits in loops
`migrate deploy` in CI	db push in prod
map P-codes centrally	try/catch everywhere

LOGGING & MONITORING

SIGNAL	HOW
Slow queries	`$on('query')` + threshold
Traces	`@prisma/instrumentation`
Errors	log `code` + `meta` + model
Pool pressure	connection timeouts
Schema drift	`migrate status` in CI

⚠ COMMON MISTAKES

- Prisma calls scattered through controllers — hard to change
- Returning Prisma models directly — leaks new columns
- One giant `schema.prisma` in big apps — merge conflicts
- Inconsistent names (`user_id` vs `authorId`) — confusing API

✓ TIPS & BEST PRACTICES

- Keep transactions in services, queries in repos
- Export `typeof db` as the app's DB type
- Use `@@map` to keep DB snake_case with camelCase code
- Review generated SQL for hot paths before launch

QUICK SUMMARY

- one client + `$extends` — single entry point
- shapes + `GetPayload` — consistent responses
- multi-file schema — organize by domain
- repo → service → route — clear layers

Security & Data Protection

SECURING DATA ACCESS

Prisma's query API and tagged-template raw queries are **parameterized**, so injection risk comes from `…Unsafe` methods and `Prisma.raw`. The bigger risks are **authorization** (users reading other users' rows) and **over-exposure** (returning secrets). Scope every query by owner or tenant, strip sensitive fields, encrypt connections, and give the app a least-privilege database user.

NOTES

- Prisma has no built-in authorization
- Postgres RLS needs per-tx `set_config`
- Encryption at rest = DB/provider setting
- Field-level crypto via an extension

SCOPING, MASKING & RLS

```ts
// tenant/owner scoping for every query
export const forTenant = (tenantId: number) =>
  prisma.$extends({ query: { $allModels: {
    async $allOperations({ args, query }) {
      if ("where" in args) args.where = { ...args.where, tenantId };
      return query(args);
    },
  } } });

// never return secrets
const prisma = new PrismaClient({
  adapter, omit: { user: { password: true, totpSecret: true } },
});

// Postgres row-level security: set the session user inside the tx
await prisma.$transaction(async (tx) => {
  await tx.$executeRaw`SELECT set_config('app.user_id', ${String(userId)}, true)`;
  return tx.invoice.findMany();           // RLS policy filters rows
});

// ✖ injectable           ✔ parameterized
prisma.$queryRawUnsafe(`SELECT * FROM users WHERE email = '${email}'`);
prisma.$queryRaw`SELECT * FROM users WHERE email = ${email}`;
```

CODE EXPLAINED

- → `$allOperations` — inject `tenantId`
- → `"where" in args` — skip creates
- → `omit` — secrets hidden by default
- → `set_config(…, true)` — tx-local setting
- → RLS policy reads `current_setting`
- → `…Unsafe` + interpolation — injectable
- → tagged template — always bound

THREAT → CONTROL

THREAT	CONTROL
SQL injection	tagged `$queryRaw`, no `Unsafe` concat
IDOR / cross-tenant	scope `where` by owner/tenant
Secret leakage	`omit`, `select` shapes
Mass assignment	Zod allow-list of fields
Sniffed traffic	`sslmode=require` / TLS
Stolen DB creds	secret manager, rotation

DATABASE HARDENING

PRACTICE	DETAIL
App DB user	DML only, no DDL
Migration user	DDL, used by CI only
Encryption at rest	provider / disk level
Field encryption	extension with app key
Backups	encrypted, access-logged
Data masking	`result` extension

⚠ COMMON MISTAKES

- `findUnique({ where: { id } })` with no owner check — IDOR
- Passing `req.body` straight into `data` — mass assignment
- Logging `query` params in production — leaks PII
- Using the migration (DDL) user at runtime

✓ TIPS & BEST PRACTICES

- Add `tenantId` to compound uniques: `@@unique([tenantId, id])`
- Hash passwords (argon2/bcrypt) before `create`
- Allow-list sortable/filterable fields from query strings
- Run `npm audit` and keep Prisma updated

QUICK SUMMARY

- parameterized queries — no string SQL
- scope by owner/tenant — authorization
- `omit` / `select` — minimal exposure
- TLS + least-privilege DB users

Migration Strategies

ZERO-DOWNTIME SCHEMA CHANGES

During a rolling deploy, old and new app versions run against the same database. Every migration must therefore be **backward compatible** with the running code. Use **expand → migrate → contract**: add new structures, backfill and dual-write, switch reads, and only then remove the old structures in a later deploy.

NOTES

- Prisma has no down migrations
- Rollback = new forward migration
- Back up before `migrate deploy`
- Renames look like drop + add — edit SQL

RENAME A COLUMN SAFELY (EXPAND → CONTRACT)

```sql
-- Deploy 1 · expand: add nullable column
ALTER TABLE "User" ADD COLUMN "fullName" TEXT;
-- app writes BOTH name and fullName, reads name

-- Deploy 2 · backfill in batches (script or SQL)
UPDATE "User" SET "fullName" = "name"
WHERE "fullName" IS NULL AND id BETWEEN 1 AND 10000;
-- app reads fullName, still writes both

-- Deploy 3 · contract: enforce + drop old column
ALTER TABLE "User" ALTER COLUMN "fullName" SET NOT NULL;
ALTER TABLE "User" DROP COLUMN "name";

-- Simple rename (needs code + DB in one step):
-- prisma migrate dev --create-only, then replace DROP/ADD with:
ALTER TABLE "User" RENAME COLUMN "name" TO "fullName";

-- Large Postgres table: build index without blocking writes
CREATE INDEX CONCURRENTLY "User_email_idx" ON "User"("email");
```

CODE EXPLAINED

- → expand — additive, old code unaffected
- → dual-write — both columns stay in sync
- → batched backfill — avoid long locks
- → contract — only after no code uses `name`
- → `--create-only` — edit SQL before apply
- → `RENAME` — keeps data (Prisma would drop)
- → `CONCURRENTLY` — keep in its own migration

CHANGE SAFETY

CHANGE	SAFE ONLINE?
Add optional column / table	✔
Add column with default	✔ (PG 11+)
Add index	✔ with CONCURRENTLY
Rename column/table	✖ expand-contract
Drop column	✖ until unused
Make column required	✖ backfill first
Change column type	✖ new column + copy

DEPLOY CHECKLIST

STEP	COMMAND / ACTION
Backup	`pg_dump` / provider snapshot
Check pending	`prisma migrate status`
Apply	`prisma migrate deploy`
Failed migration	fix SQL → `migrate resolve`
Roll back	new forward migration
Verify	smoke tests, error rates

⚠ COMMON MISTAKES

- Renaming a field in the schema — Prisma drops and re-adds the column
- Dropping a column in the same deploy that stops using it
- One huge `UPDATE` backfill — long locks, replication lag
- Editing applied migrations to "fix" production

✓ TIPS & BEST PRACTICES

- Practice the deploy on a copy of production data
- Set a `lock_timeout` in risky migration SQL
- Keep contract steps in a separate, later release
- Automate `migrate deploy` in CI before app rollout

QUICK SUMMARY

- expand → migrate → contract — zero downtime
- `--create-only` — hand-edit SQL
- backups + `migrate status` — before deploy
- rollback — forward-fix migration

Sheet 33

Quick Reference — Schema & CLI

CLI commands · prisma.config.ts · generator & datasource · field types · attributes · relations · enums · migrations workflow

CLI

COMMAND	DOES	#
`prisma init`	schema + config scaffold	01
`prisma generate`	build the client	02
`prisma migrate dev`	create + apply migration	17
`prisma migrate deploy`	apply in prod / CI	17
`prisma db push`	sync schema, no migration	17
`prisma db pull`	introspect existing DB	32
`prisma studio`	data browser GUI	25

SETUP (PRISMA 7)

ITEM	VALUE	#
`prisma.config.ts`	schema path, migrations, datasource URL	01
`generator client`	`provider = "prisma-client"`	02
`output`	required, e.g. `"../src/generated/prisma"`	02
`datasource db`	`provider = "postgresql"`	01
`@prisma/adapter-pg`	driver adapter for Postgres	20
`new PrismaClient({ adapter })`	create the client	02

FIELD TYPES

TYPE	DB MAPPING	#
`String`	text / varchar	03
`Int / BigInt`	integer / bigint	03
`Float / Decimal`	double / numeric	03
`Boolean`	boolean	03
`DateTime`	timestamp	03
`Type? / Type[]`	optional / list	03

FIELD ATTRIBUTES

ATTRIBUTE	DOES	#
`@id`	primary key	05
`@default(autoincrement())`	serial int	05
`@default(uuid()) / cuid()`	generated id	05
`@default(now())`	timestamp default	05
`@updatedAt`	set on every update	05
`@unique`	unique constraint	05
`@map("col") / @db.VarChar(80)`	column name / native type	05

MODEL ATTRIBUTES

ATTRIBUTE	DOES	#
`@@id([a, b])`	composite key	05
`@@unique([email, tenantId])`	composite unique	05
`@@index([createdAt])`	index	19
`@@map("users")`	table name	05

RELATIONS

SYNTAX	DOES	#
`posts Post[]`	one-to-many side	04
`author User @relation(fields: [authorId], references: [id])`	FK side	04
`authorId Int`	scalar FK field	04
`onDelete: Cascade`	referential action	04
`profile Profile?`	one-to-one	04
`tags Tag[] (both sides)`	implicit many-to-many	14
`@relation("Mentor")`	disambiguate / self relation	28

MIGRATIONS

TASK	COMMAND / NOTE	#
`New migration`	`migrate dev --name add_user`	17
`Draft only`	`migrate dev --create-only`	17
`Reset dev DB`	`migrate reset` (drops data)	17
`Mark applied`	`migrate resolve --applied <name>`	32
`Check drift`	`migrate diff`	32
`Baseline existing DB`	`db pull` + diff + resolve	32

⚠ COMMON MISTAKES

- Forgetting `prisma generate` after a schema change — the client types go stale
- Missing `output` in the Prisma 7 generator, or importing from `@prisma/client`
- Using `db push` in production — it skips migration history

✓ TIPS & BEST PRACTICES

- Index every foreign key you filter on (`@@index([authorId])`)
- Review generated SQL with `--create-only` before applying risky changes
- Keep the connection URL in `prisma.config.ts` / env, never in the schema

QUICK SUMMARY

- Schema = generator + datasource + models
- `migrate dev` locally, `migrate deploy` in prod
- `@id`, `@unique`, `@default`, `@relation` do most of the work
- Prisma 7: `prisma-client` generator + driver adapter

Sheet 34

Quick Reference — Client Queries

CRUD methods · filters & operators · relation queries · nested writes · pagination · aggregation · transactions · raw SQL · errors

CRUD METHODS

METHOD	DOES	#
`create({ data })`	insert one	06
`createMany / createManyAndReturn`	insert many	06
`findUnique({ where })`	by unique field	07
`findFirst / findMany`	filter + sort	07
`findUniqueOrThrow`	throw if missing	07
`update / updateMany`	change rows	08
`upsert({ where, create, update })`	insert or update	08
`delete / deleteMany`	remove	09

FILTERS

OPERATOR	EXAMPLE	#
`equals / not`	`{ name: { not: "x" } }`	10
`in / notIn`	`{ id: { in: [1, 2] } }`	10
`lt / lte / gt / gte`	`{ age: { gte: 18 } }`	10
`contains / startsWith`	`{ email: { endsWith: ".io" } }`	10
`mode: "insensitive"`	case-insensitive (PG)	10
`AND / OR / NOT`	`{ OR: [{…}, {…}] }`	11
`some / every / none`	filter on relation lists	11

SHAPE THE RESULT

OPTION	DOES	#
`select: { id: true }`	pick fields	12
`include: { posts: true }`	load relations	12
`omit: { password: true }`	drop fields	12
`orderBy: { createdAt: "desc" }`	sort	07
`distinct: ["city"]`	unique by field	07
`_count: { select: { posts: true } }`	count relations	12

NESTED WRITES

OPERATION	DOES	#
`create`	create related rows	13
`connect: { id }`	link existing	13
`connectOrCreate`	link or create	13
`disconnect / set`	unlink / replace list	13
`update / upsert (nested)`	change related	13
`delete (nested)`	remove related	13

PAGINATION

OPTION	DOES	#
`take: 20`	limit	11
`skip: 40`	offset	11
`cursor: { id: lastId }`	cursor paging	11
`skip: 1 + cursor`	skip the cursor row	11
`take: -20`	page backwards	11

AGGREGATION

METHOD	DOES	#
`count({ where })`	number of rows	15
`aggregate({ _sum, _avg })`	sums, averages	15
`_min / _max`	extremes	15
`groupBy({ by: ["role"] })`	group rows	15
`having: { … }`	filter groups	15

TRANSACTIONS & RAW

API	DOES	#
`$transaction([a, b])`	batch, all or nothing	16
`$transaction(async tx => …)`	interactive	16
`isolationLevel / timeout`	tx options	16
`$queryRaw`SELECT …${x}``	parameterized query	26
`$executeRaw`…``	returns row count	26
`$queryRawUnsafe`	dynamic SQL — avoid	31
`$extends({ model: … })`	custom methods	22

ERRORS

CODE	MEANS	#
`P2002`	unique constraint failed	18
`P2025`	record not found	18
`P2003`	foreign key failed	18
`P2034`	transaction conflict, retry	18
`PrismaClientKnownRequestError`	check `e.code`	18
`PrismaClientValidationError`	bad query shape	18

⚠ COMMON MISTAKES

- N+1: querying relations in a loop — use `include` or one `findMany` with `in`
- `$queryRawUnsafe` with string-built SQL (injection)
- `deleteMany({})` with an empty filter wipes the table

✓ TIPS & BEST PRACTICES

- Use `select` to return only the fields the caller needs
- Prefer cursor pagination for large or live tables
- Turn on query logging (`log: ["query"]`) while tuning

QUICK SUMMARY

- `findMany({ where, select, include, orderBy, take })`
- Nested writes: `create` , `connect` , `connectOrCreate`
- `$transaction` for all-or-nothing
- Catch `P2002` / `P2025` by `e.code`